

Multi-channel Advertising

Four methods for measuring advertising response

Dr. Junbum Kwon

29 March 2019

Author: Elea McDonnell Feit, Professor at Drexel University and Senior Fellow at Wharton Customer Analytics Initiative

Topics

Why advertising response?

Attribution rules

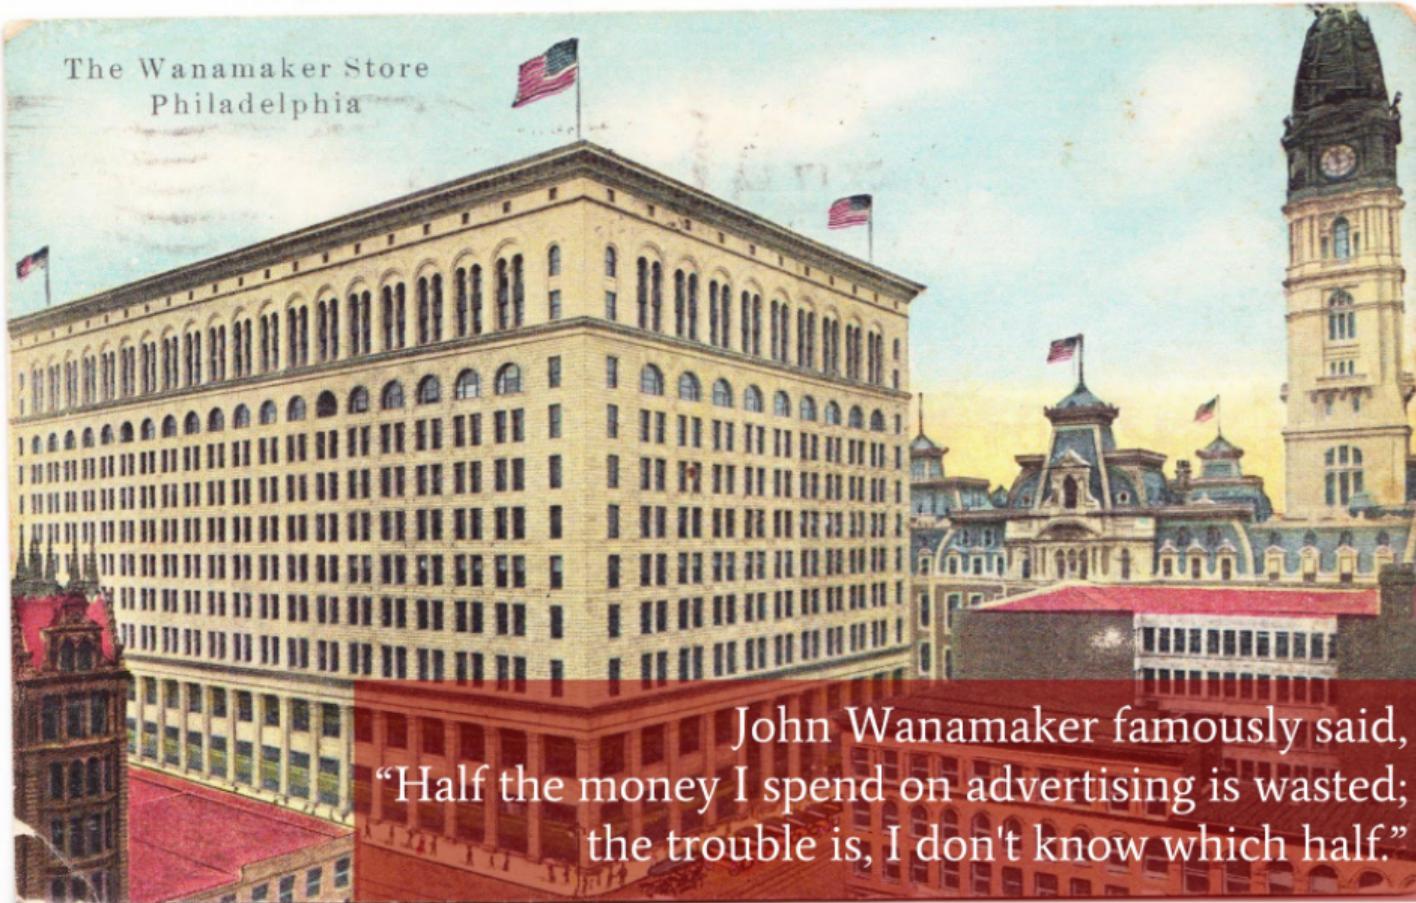
Holdout testing

Marketing mix modeling

Model-based attribution

Why advertising response?

The Wanamaker Store
Philadelphia



John Wanamaker famously said,
“Half the money I spend on advertising is wasted;
the trouble is, I don't know which half.”

Measuring advertising response

The goal of any marketing campaign is to increase sales (short-term or long-term).

Measuring advertising response

The goal of any marketing campaign is to increase sales (short-term or long-term).

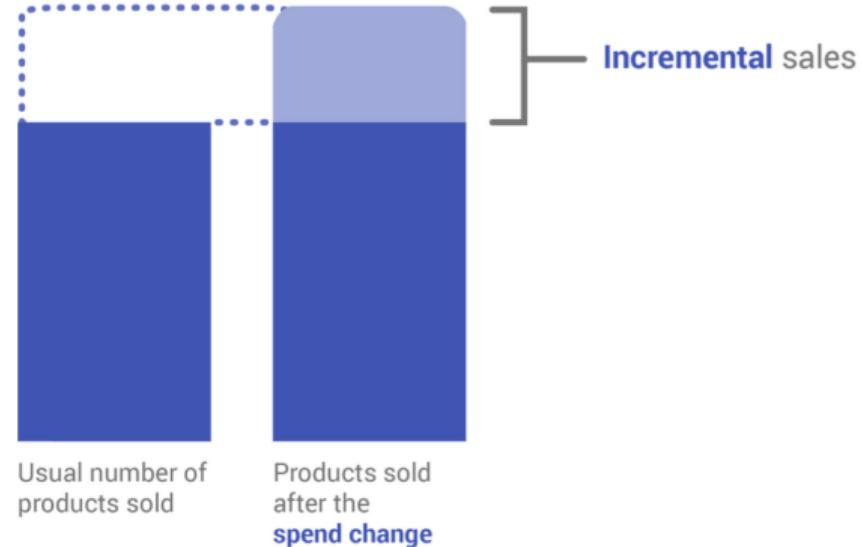
In theory, it should be easy to evaluate the performance of marketing. Each campaign or marketing channel should be evaluated based on the **incremental profit** that it produces relative to its **cost**.

$$ROI = \frac{\text{incremental profit due to advertising} - \text{cost of advertising}}{\text{cost of advertising}}$$

ROI = Return on Investment

Incremental sales

Incremental profit depends on **incremental sales** which are the additional sales we make with advertising over and above what we would have sold without advertising. Incremental profit is typically a function of the incremental sales.

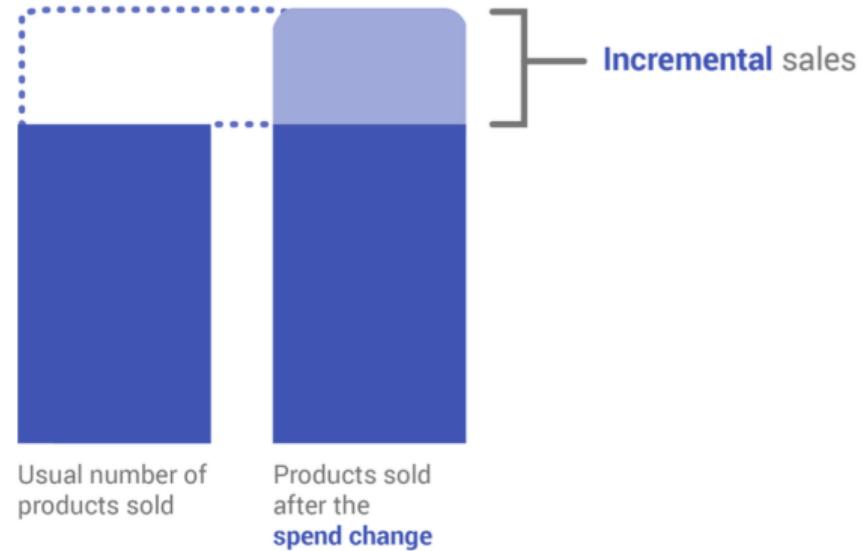


Source: [Think with Google](#)

Incremental sales

Incremental profit depends on **incremental sales** which are the additional sales we make with advertising over and above what we would have sold without advertising. Incremental profit is typically a function of the incremental sales.

With today's digital media – where we can track individual users – it is now possible to estimate incremental sales in many situations.



Source: [Think with Google](#)

Goals

Understand four popular techniques for estimating the incremental sales due to advertising.

- Attribution rules
- Holdout testing
- Marketing mix modeling (with aggregate spending and sales data)
- Attribution modeling (with user-level data)

Learn how each technique works using the Python.

Understand which techniques produce better estimates of incremental profits and why.

Why advertising response?

Data

Data: retail impressions and conversions

We will be analyzing a data set describing 10,000 customers and potential customers of a retailer. The retailer uses four different advertising channels.

Display

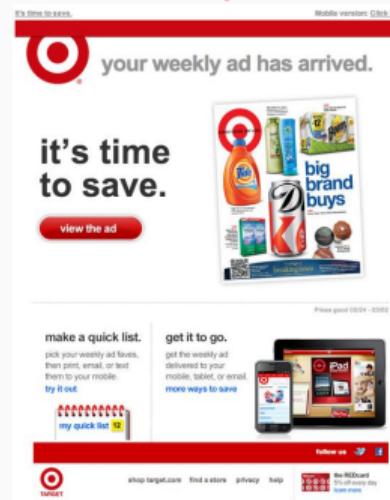


Social

Sponsored

A social media post from Target's official account. It shows a woman wearing a green and white floral skirt and a tan wedge sandal. The post is labeled 'Sponsored' at the top left. On the right side, there is a 'Create Ad' button. Below the image, the text reads 'Target Style target.com Love these? Make 2017 your most stylish year yet.'

Email



Direct (via mail)



Customer tracking begins when the customer is exposed to a display or social ad, visits the retailer website or makes a purchase.

Data structure

The data is organized in three files:

- `customer.csv`: each row is a customer, 10,000 rows
- `impressions.csv`: each row is an exposure of a marketing communication to a specific customer, 501,336 rows
- `transactions.csv`: each row is a transaction made by a customer

Data structure

The data is organized in three files:

- `customer.csv`: each row is a customer, 10,000 rows
- `impressions.csv`: each row is an exposure of a marketing communication to a specific customer, 501,336 rows
- `transactions.csv`: each row is a transaction made by a customer

This basic structure is typical of raw digital advertising data. The data is hypothetical, and has been designed to illustrate important points in data analysis.

Customer file

The customer file can be downloaded from goo.gl/fyxRE5. Each row describes a customer.

Columns:

id: an id number for the customer

past purchase: whether the customer has made a purchase prior to the observation period.

email: indicates whether the customer is eligible to receive emails, i.e. we have an email address and permission to mail

direct: indicates whether the customer is eligible to receive direct mail, i.e. we have an address and permission to mail

Read and inspect customer file

We can read the file directly into R:

```
> cust <- read.csv("https://goo.gl/mqy8NR")
> nrow(cust)
[1] 10000
> summary(cust)
```

	<code>id</code>	<code>past.purchase</code>	<code>email</code>	<code>direct</code>
Min.	: 1	Min. :0.0000	Min. :0.0000	Min. :0.0000
1st Qu.	: 2501	1st Qu.:0.0000	1st Qu.:0.0000	1st Qu.:0.0000
Median	: 5000	Median :1.0000	Median :1.0000	Median :0.0000
Mean	: 5000	Mean :0.5022	Mean :0.6001	Mean :0.4974
3rd Qu.	: 7500	3rd Qu.:1.0000	3rd Qu.:1.0000	3rd Qu.:1.0000
Max.	:10000	Max. :1.0000	Max. :1.0000	Max. :1.0000

We have 10,000 customers and about half of those have made a purchase. About 60% are eligible for email and 50% are eligible for direct mail.

More inspection of customer file

```
> xtabs(~email+direct+past.purchase , data=cust)
, , past.purchase = 0
    direct
email      0      1
  0 2775  707
  1 1205  291
, , past.purchase = 1
    direct
email      0      1
  0    95   422
  1  951  3554
```

Most customers who have not made a purchase are not eligible for email or direct mail.
Customers who have made a purchase are more likely to be eligible.

Impressions file

The impressions file can be downloaded from <https://goo.gl/LT93AA>. Each row in the file represents an exposure of one customer to an ad, i.e. an impression.

Columns:

id: id number for the customer

date: date of impression. Most files would have a date-time, but we have simplified for the workshop.

channel: channel of the ad exposure

click: indicates whether the customer clicked on the ad

Read and inspect **impressions** file

```
> impress <- read.csv("https://goo.gl/74qIxY")
> impress$date <- as.Date(impress$date) # change type
> nrow(impress)
[1] 501336
> summary(impress)
      id              date            channel        click
 Min.   : 1   Min.   :2016-12-31   direct     : 9948   Min.   :0.00000
 1st Qu.: 2467  1st Qu.:2017-01-10  display    :216371  1st Qu.:0.00000
 Median : 4940  Median :2017-01-20  email     : 38426  Median :0.00000
 Mean   : 4960  Mean   :2017-01-22  email.holdout: 9582   Mean   :0.01854
 3rd Qu.: 7454  3rd Qu.:2017-01-31  social    :227009  3rd Qu.:0.00000
 Max.   :10000  Max.   :2017-02-27
```

We have 501,336 impressions between 2016-12-31 to 2017-02-27. The majority of these impressions are social or display. Overall click rate is 1.85%, which is reasonable.

Summarize the cadence of the impressions

Cadence is the timing of advertising impressions. It is useful to understand the cadence of your ads before you do any other analysis. Email and direct are often sent out to users on specific dates. Display and social ads can be steady or can be turned on and off at specific times. We can summarize the cadence with the crosstab function in R.

```
> (cadence <- xtabs(~date+channel, data=impress))
```

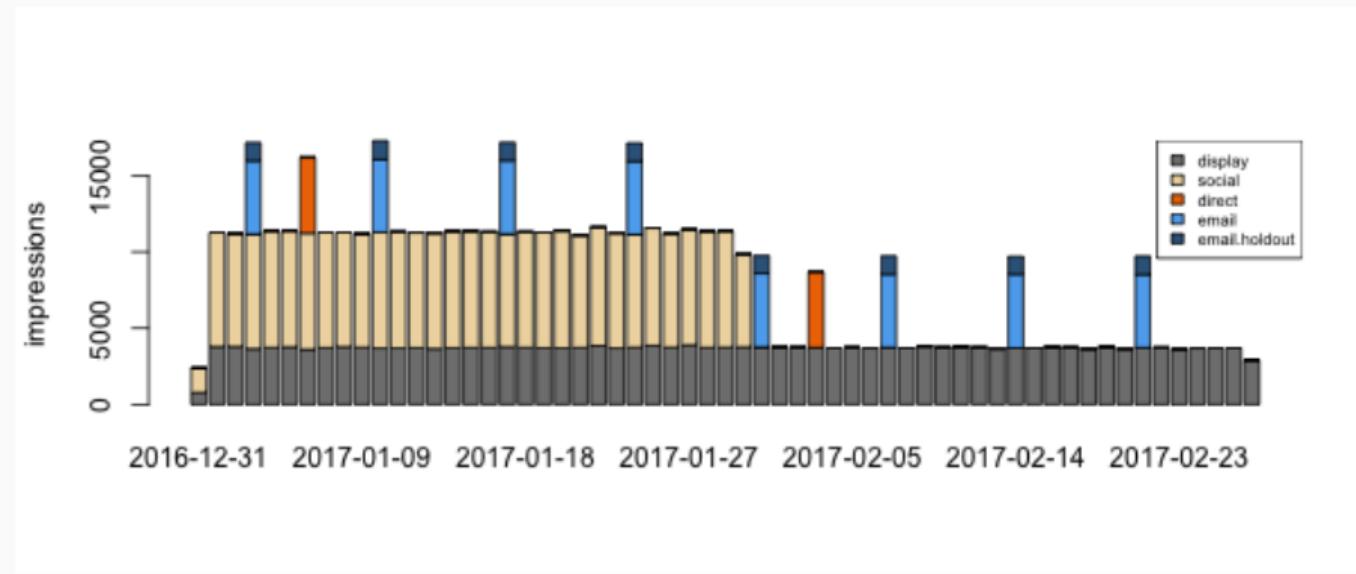
date	channel				
	direct	display	email	email.holdout	social
2016-12-31	0	788	0	0	1610
2017-01-01	0	3786	0	0	7481
2017-01-02	0	3792	0	0	7416
2017-01-03	0	3656	4798	1203	7505
2017-01-04	0	3731	0	0	7648
...					

Visualizing the cadence

It is useful to create plot of the cadence data.

```
> cadence <- cadence[,c(2,5,1,3,4)] # reorder the columns  
> mycols <- c("gray50", "wheat2", "darkorange2", "steelblue2", "steelblue4")  
> barplot(t(cadence), col=mycols, ylab="impressions")  
> legend("topright", legend=colnames(cadence), fill=mycols, cex=0.6)
```

Visualizing the cadence



Display impressions per day are steady across the observation window.

Social is steady in the first month and then stops.

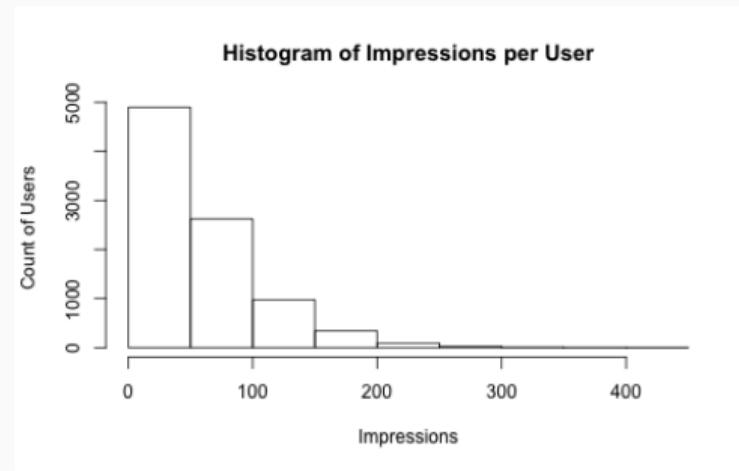
Emails are sent once per week and each campaign seems to have a holdout.

Direct mail is sent once per month.

Histogram of impressions

It is also useful to understand how many impressions each customer gets.

```
> hist(xtabs(~id, data=impress), xlab="Impressions", ylab="Count of Users",  
+       main="Histogram of Impressions per User")
```



Some customers receive as many as 400 impressions in two months, but most get less than 100 impressions.

Click through rates by channel

```
> xtabs(click~channel, data=impress)/xtabs(~channel, data=impress)
channel
    direct      display      email  email.holdout      social
0.0000000000  0.004783451  0.099672097  0.000000000  0.019519050
```

The click through rates are highest for email at 10.0% and lowest for display at 0.5%. As expected, there are no clicks for direct or email holdouts.

Transactions file

The transactions file can be downloaded from <https://goo.gl/ZErZ6g>. Each row in the file represents a purchase made by a customer.

Columns:

`id`: id number for the customer

`date`: date of transactions. Most files would have a date-time, but we have simplified for the workshop.

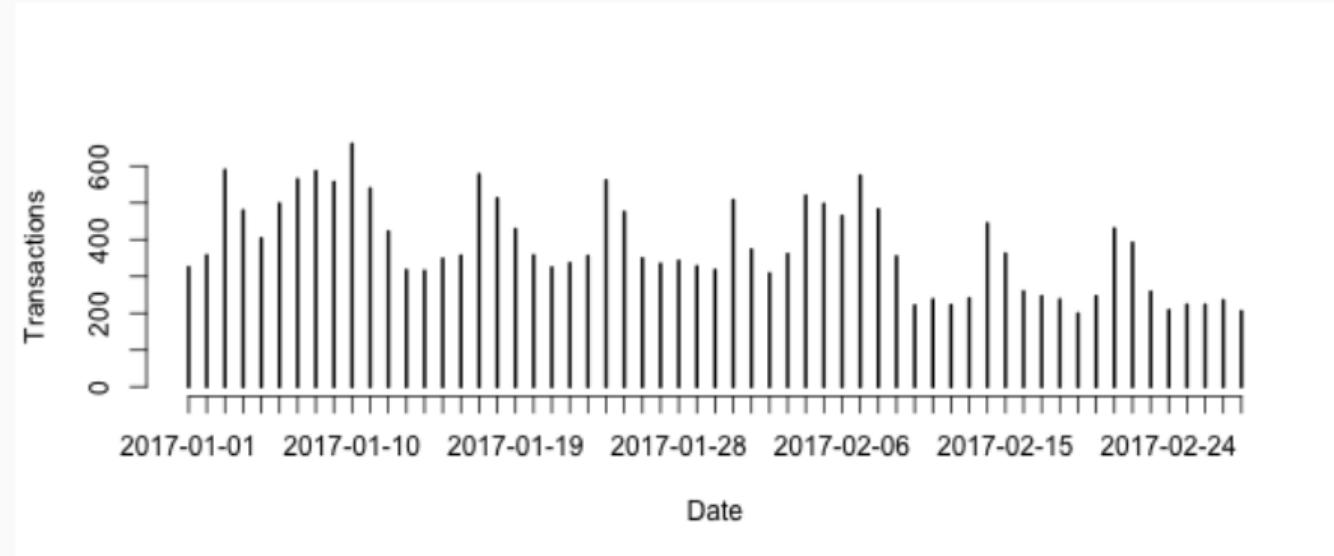
`last.touch`: channel of the last ad impression the customer saw before the transaction

`last.click`: channel of the last ad the customer clicked before the transaction

Purchase amounts can be used in holdout testing, attribution models or marketing mix models, but we will focus on transaction counts for this workshop.

Plot of transactions over time

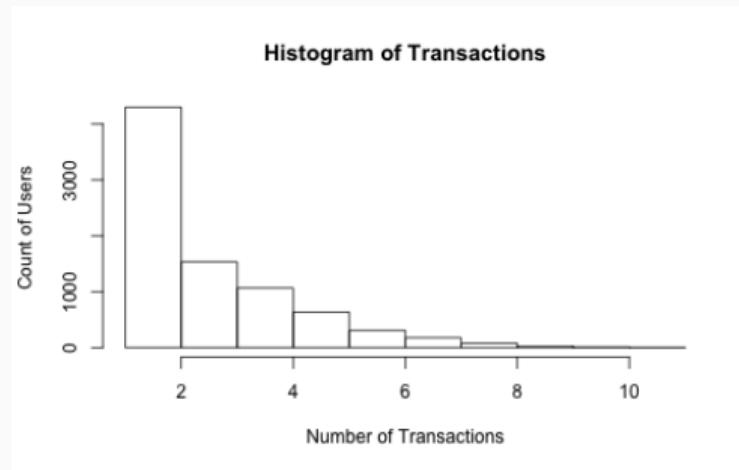
```
> plot(transbyday, ylab="Transactions", xlab="Date")
```



Transactions appear to be a bit higher in the first month and there are clear spikes around the time of emails and direct mail.

Histogram of transactions

```
> hist(xtabs(~id, data=trans), xlab="Number of Transactions", ylab="Count of Users",
+       main="Histogram of Transactions")
```



Most customers have less than 2 transactions in the two month observation period.

Customer 100

To understand how the files fit together, it is helpful to look at the impressions and transactions for one customer.

```
> cust[cust$id==100,]
  id past.purchase email direct
100 100            0    0    0
> impress[impress$id==100,]
[1] id      date    channel click
<0 rows> (or 0-length row.names)
> trans[trans$id==100,]
  X id      date last.touch last.click
206 206 100 2017-01-18      none      none
207 207 100 2017-01-26      none      none
```

Customer 100 has no impressions and made two transactions.

Customer 300

```
> cust[cust$id==300,]
  id past.purchase email direct
300 300           1     1     1
> summary(impress[impress$id==300,])
      id          date            channel      click
Min. :300  Min.   :2016-12-31  direct       : 2  Min.   :0
1st Qu.:300  1st Qu.:2017-01-10  display      :81  1st Qu.:0
Median :300  Median  :2017-01-29  email        : 8  Median  :0
Mean   :300  Mean    :2017-01-28  email.holdout: 0  Mean   :0
3rd Qu.:300  3rd Qu.:2017-02-14  social       : 0  3rd Qu.:0
Max.   :300  Max.   :2017-02-26                    Max.   :0
> trans[trans$id==300,]
      X id      date last.touch last.click
647 647 300 2017-01-03   display      none
648 648 300 2017-01-10   email        none
649 649 300 2017-02-04   direct       none
650 650 300 2017-02-19   display      none
```

Customer 300 has 8 email, 2 direct mail and 81 display impressions and made 4 transactions.

Steps in data inspection

1. For each file,
 - 1.1 Make sure you have the entire file read in by checking number of rows.
 - 1.2 Summarize individual variables in file mean, min, max, etc. Do they make sense?
 - 1.3 Summarize relationships between variables in file using crosstabs or scatterplots. Do they make sense?
2. Check that joins between files are as expected (skipped).
3. Summarize relationships between files. Do they make sense?

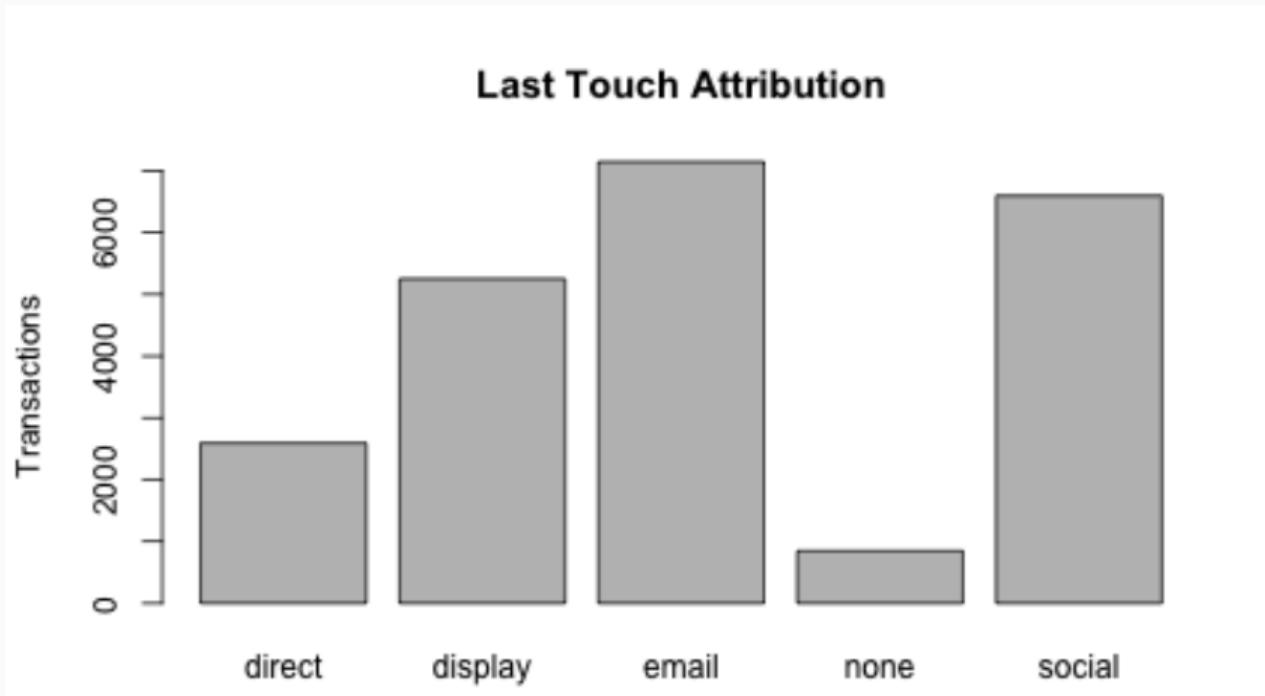
Attribution rules

Attribution rules

What is last click?

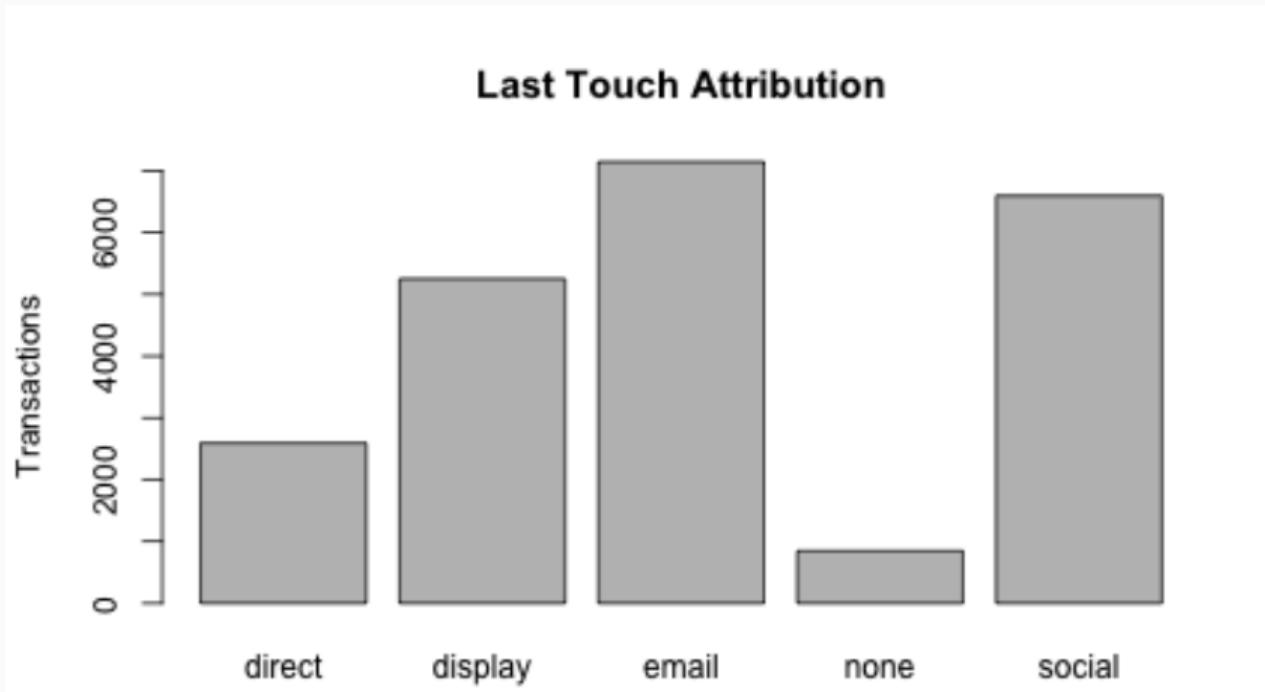
Last-touch attribution

Last-click (or last-touch) attribution looks backward from each conversion to find the last ad the user clicked on prior to the conversion.



Last-touch attribution

Last-click (or last-touch) attribution looks backward from each conversion to find the last ad the user clicked on prior to the conversion.



Attribution rules

Last touch analysis

Last touch analysis

The last click and the last touch are stored in the transaction file.

```
> head(trans)
  id      date last.touch last.click
1  2 2017-01-04      email      none
2  2 2017-02-12      email      none
3  3 2017-02-02      email      none
4  3 2017-02-14      email      none
5  5 2017-01-04    display      email
6  5 2017-01-13    display      email
```

One advantage of last-click is that it is easy to compute once this pre-processing is done.

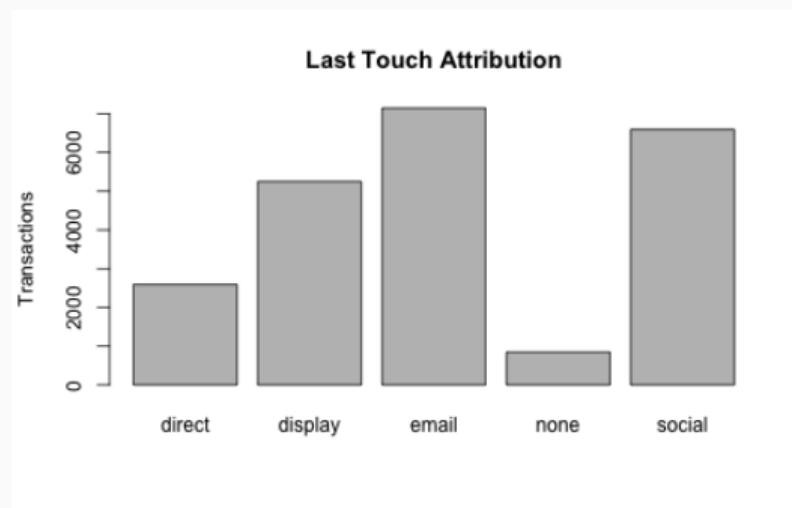
Last touch analysis

For example, if we want to compute the number of transactions that are attributed to each channel by last touch, we just do a quick crosstab on the transaction table.

```
> (last.touch.att <- xtabs(~last.touch, data=trans))
last.touch
direct display   email     none   social
    2594      5252     7145      846     6596
> barplot(last.touch.att, ylab="Transactions",
+           main="Last Touch Attribution")
```

When we do this, we are ignoring all the customers who didn't transact.

Last touch analysis



Many people interpret this as meaning that the incremental sales for social are 6596.

Last touch for subgroups of transactions

It is also easy to compute the last touch for a subgroup of the transactions. Simply subset out the transactions of interest and then crosstab.

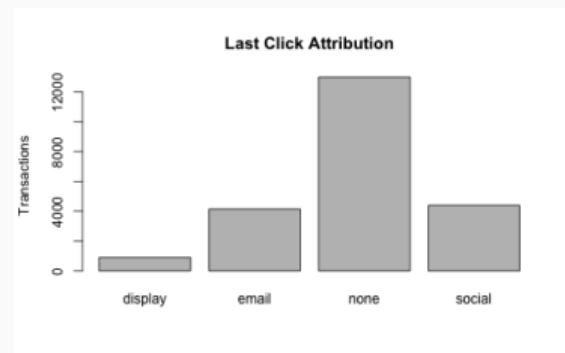
```
> xtabs(~last.touch, data=trans[trans$date>as.Date("2017-01-31"),])  
last.touch  
direct display email none social  
1669    2682    4081    258    328
```

In February, there were far fewer sales attributed to social, most likely because we ended the social ads at the end of January.

Last click analysis

You can do the same analysis based on the last ad a customer clicked, rather than the last touch/impression.

```
> (last.click.att <- xtabs(~last.click, data=trans))
last.click
display   email     none   social
      891     4151    12998    4393
> barplot(last.click.att, ylab="Transactions",
+           main="Last Click Attribution")
```

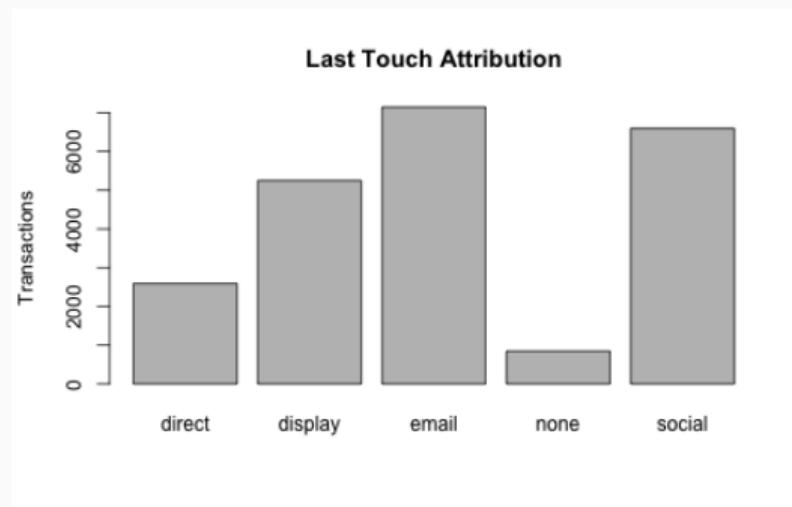


Since people don't click very much, the attributed sales are much lower. Measuring ad performance based on clicks doesn't make much sense for the advertiser.

Attribution rules

Limitations of last-click

Last touch analysis



Many people interpret this as meaning that the incremental sales for social are 6596.

Incorrect assumptions behind last touch

When we use last touch to estimate the incremental sales for an ad, we are making two mistakes:

1. Other ads may have influenced the customer and contributed to the sale.

Incorrect assumptions behind last touch

When we use last touch to estimate the incremental sales for an ad, we are making two mistakes:

1. Other ads may have influenced the customer and contributed to the sale.
 - This is the problem everyone seems to recognize.

Incorrect assumptions behind last touch

When we use last touch to estimate the incremental sales for an ad, we are making two mistakes:

1. Other ads may have influenced the customer and contributed to the sale.
 - This is the problem everyone seems to recognize.
2. It counts all the sales as incremental, i.e. it assumes that customers who saw ads would not have bought if they hadn't seen the ads. (Recall treatment vs control group)

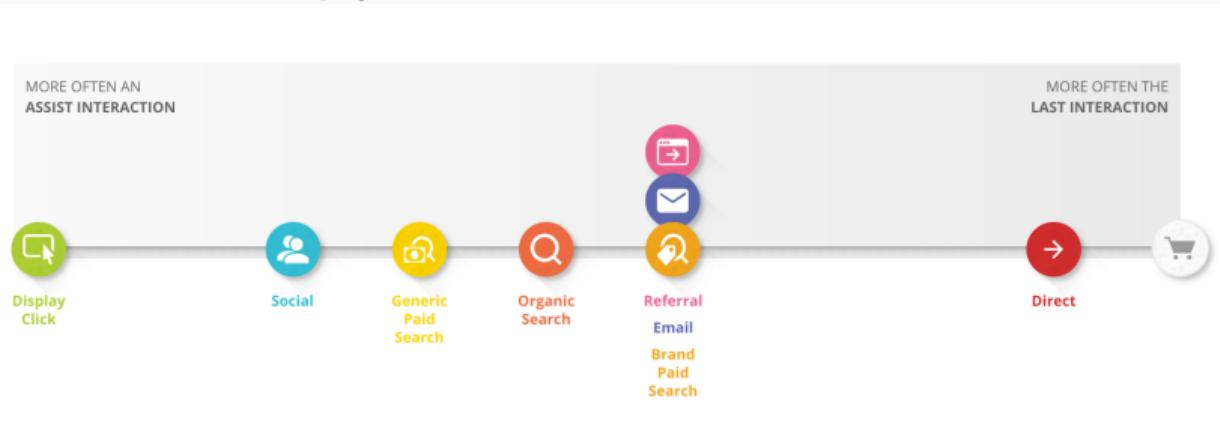
Incorrect assumptions behind last touch

When we use last touch to estimate the incremental sales for an ad, we are making two mistakes:

1. Other ads may have influenced the customer and contributed to the sale.
 - This is the problem everyone seems to recognize.
2. It counts all the sales as incremental, i.e. it assumes that customers who saw ads would not have bought if they hadn't seen the ads. (Recall treatment vs control group)
 - This issue is less well recognized.

Attribution as a payment mechanism

Because last-touch is so simple to compute, it has been used as a rule for allocating payments between channels on pay-for-conversion ads.

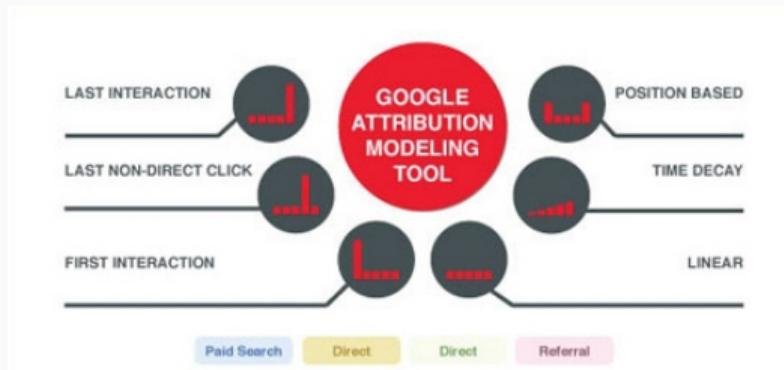


Source: [Think with Google](#)

Last touch unfairly favors channels that tend to show ads towards the end of the path to purchase such as search and retargeting.

Other attribution rules (“models”)

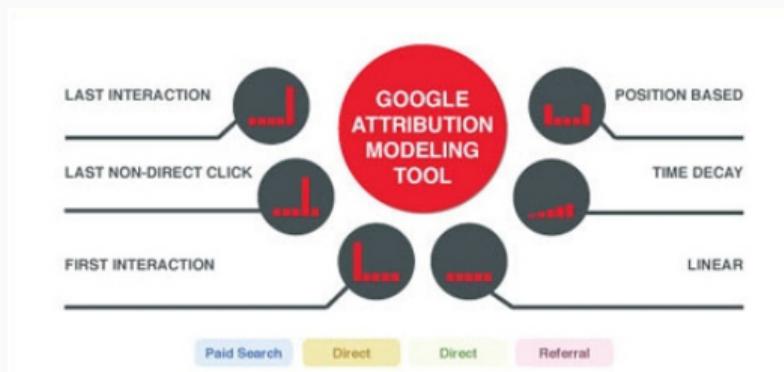
Perhaps because everyone intuitively understood that last touch didn't account for “assists” and was unfair to channels that tend to occur early in the path to purchase they came up with other arbitrary rules.



Source: [Across Health](#)

Other attribution rules (“models”)

Perhaps because everyone intuitively understood that last touch didn't account for “assists” and was unfair to channels that tend to occur early in the path to purchase they came up with other arbitrary rules.



Source: [Across Health](#)

These improved rules may help problem 1, but they don't help with problem 2.

Last touch is largely dead

There are many vendors such as Convertro, VisualIQ and MarketShare that have provided model-based attribution solutions.



Source: [Google Analytics 360](#)

In Google Analytics 360, released in 2016, Google also abandoned rule-based attribution in favor of model-based approaches.

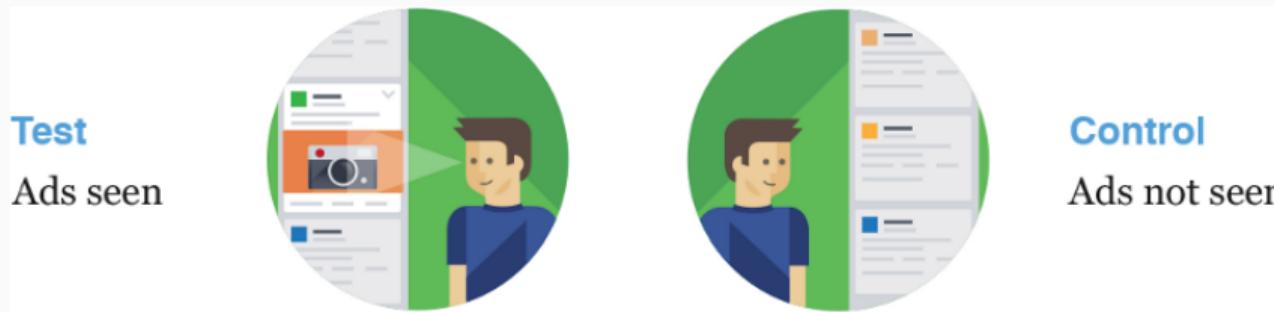
Holdout testing

Holdout testing

What is a holdout test?

Holdout testing

The gold standard for measuring incremental sales is an experiment, where we **randomly** assign customers to be exposed or not exposed to an ad.



Source: [Designing with Science, medium.com](https://designingwithscience.com)

For example, with email we can take our list of target emails and randomly select a set of them to not receive an email. Facebook also has a tool for randomized holdouts (for large advertisers.)

Why holdout testing works

The magic behind a holdout test is the randomization. Another name for holdout testing is **randomized controlled trial**.

Why holdout testing works

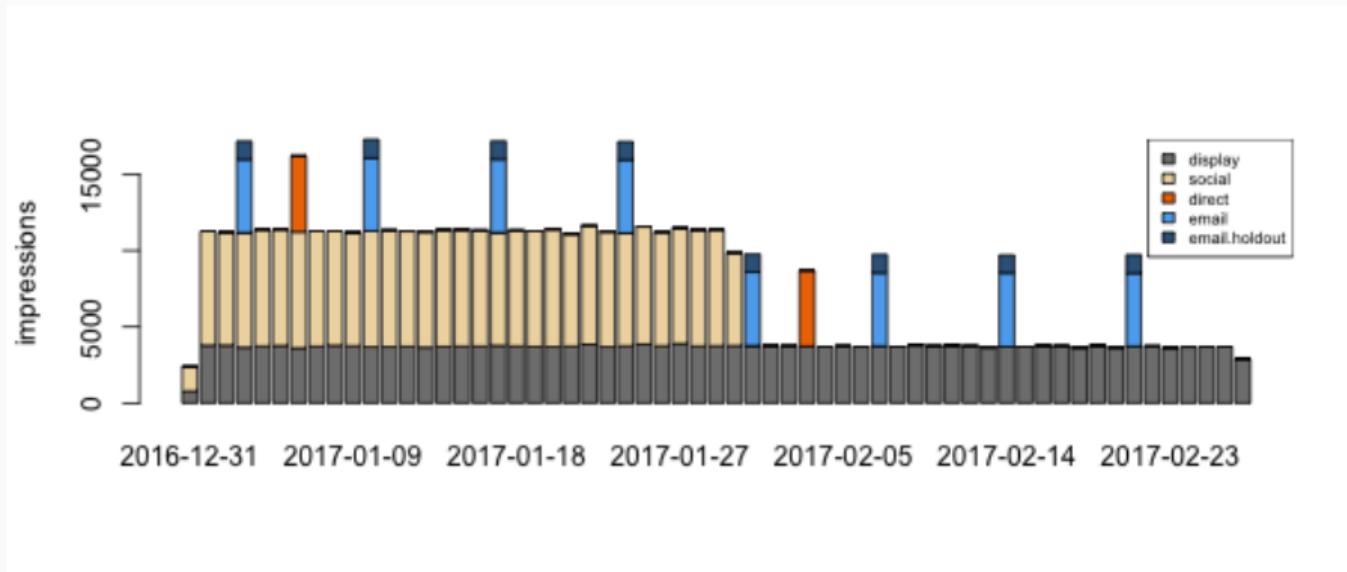
The magic behind a holdout test is the randomization. Another name for holdout testing is **randomized controlled trial**.

By randomly selecting the customers for the control group, we assure that the two groups are the same on average. The treatment and control groups will be similar in their propensity to purchase and response to ads. Statisticians call this **probabilistically equivalent**.

Holdout testing

Holdout test analysis

Email holdout tests in the example data



There was an email on 2017-01-03 that included a holdout group. Let's analyze this test.

Function for analyzing holdout tests

```
holdout.test <- function(test.date, delay=0, window, impress, trans) {  
  test.ids <- unique(impress$id[impress$channel=="email" & impress$date==test.date])  
  control.ids <- unique(impress$id[impress$channel=="email.holdout" & impress$date==test.date])  
  tdata <- data.frame(ids=c(test.ids, control.ids))  
  tdata$group[tdata$id %in% test.ids] <- "test"  
  tdata$group[tdata$id %in% control.ids] <- "control"  
  in.window <- trans$date>=(test.date+delay) & trans$date<(test.date+window+delay)  
  tdata$convert <- tdata$id %in% trans$id[in.window]  
  ttable <- xtabs(~ group + convert, data=tdata)  
  mosaicplot(~ group + convert, data=tdata,  
             main=paste("Email Test on", test.date))  
  proptest <- prop.test(x=ttable[, "TRUE"], n=xtabs(~ group, data=tdata))  
  diff.conv <- c(diff=proptest$estimate[2]-proptest$estimate[1], ci=-proptest$conf.int)  
  out <- list(diff.conv, ttable, proptest)  
}
```

Analysis of email holdout test

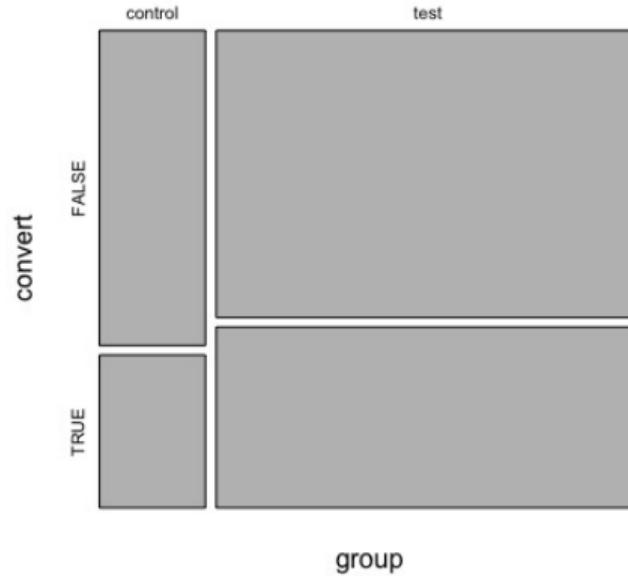
```
> (holdout.test(test.date=as.Date("2017-01-03"), window=7, impress=impress,
+                trans=trans))
[[2]]
    convert
group      FALSE  TRUE
control     810   393
test        2944  1854

[[3]]
2-sample test for equality of proportions with
continuity correction

data: ttable[, "TRUE"] out of xtabs(~group, data = tdata)
X-squared = 14.395, df = 1, p-value = 0.0001482
alternative hypothesis: two.sided
95 percent confidence interval:
 -0.09011755 -0.02933788
sample estimates:
prop 1    prop 2
0.3266833 0.3864110
```

Reporting of email holdout test

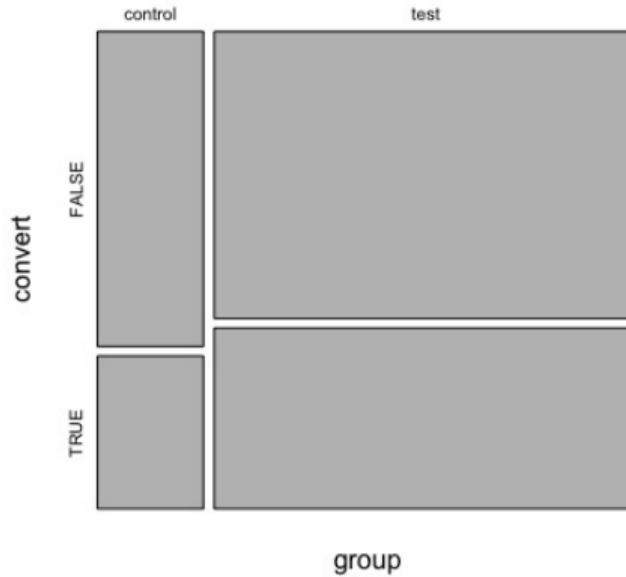
Email Test on 2017-01-03



The test group had a 38.6% conversion rate in the 7 days after the email was sent, versus a 32.7% conversion rate for the control group.

Reporting of email holdout test

Email Test on 2017-01-03



The test group had a 38.6% conversion rate in the 7 days after the email was sent, versus a 32.7% conversion rate for the control group.

The email on 2017-01-03 produced incremental sales. The incremental increase in conversion rate is between 5.9%.

Another analysis of email holdout test on 2017-01-01

```
> (holdout.test(test.date=as.Date("2017-01-03"), window=3, impress=impress, trans=trans))
[[2]]
      convert
group      FALSE  TRUE
control    1063   140
test       3850   948

[[3]]
 2-sample test for equality of proportions with continuity correction

data: ttable[, "TRUE"] out of xtabs(~group, data = tdata)
X-squared = 42.187, df = 1, p-value = 0.0000000008295
alternative hypothesis: two.sided
95 percent confidence interval:
 -0.10306428 -0.05934891
sample estimates:
 prop 1     prop 2 
0.1163757 0.1975823
```

Setting the response window

If you change the response window, you will get a different answer about how much advertising increases sales.

Setting the response window

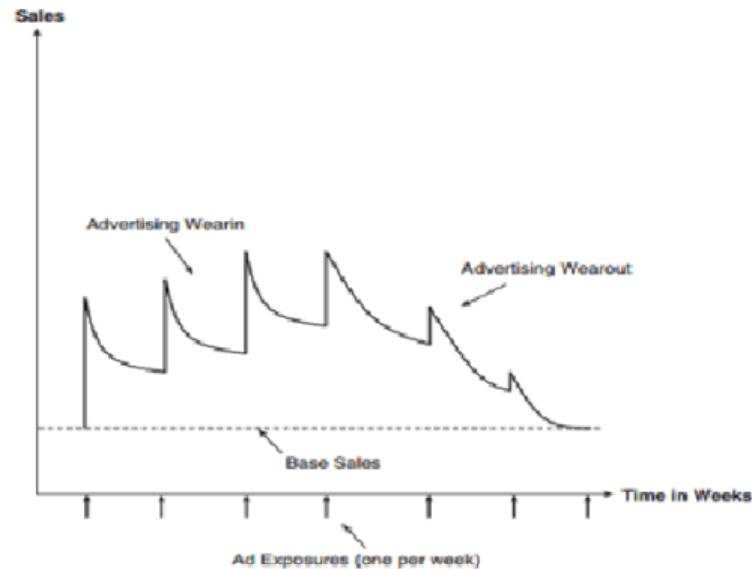
If you change the response window, you will get a different answer about how much advertising increases sales.

- For the email test on 2017-01-03, a response window of 3 days indicated a increase in conversion of 8.1% versus 6.0% for a 7 day window.

What is happening here?

Checking ad response over time

Ad response is often greatest just after exposure and then falls off over time.



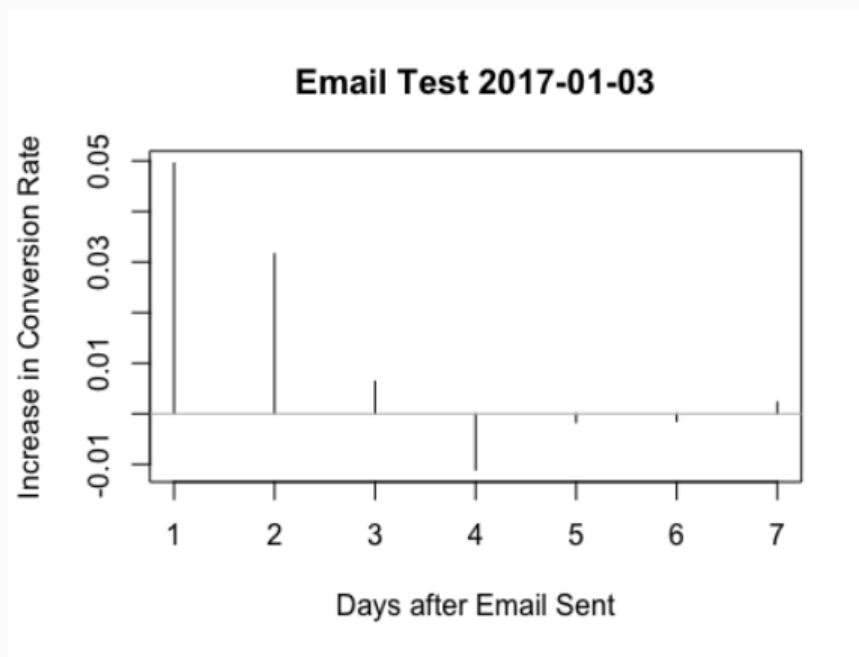
Source: [Michael Wolfe](#)

Ad response over time for 2017-01-03 email test

With a holdout test, you can study the ad response over time to learn what the pattern looks like over time.

```
day1 <- holdout.test(test.date=as.Date("2017-01-03"), delay=0, window=1,
impress=impress, trans=trans)
...
day7 <- holdout.test(test.date=as.Date("2017-01-03"), delay=6, window=1,
impress=impress, trans=trans)
incr.conv <- c(day1[[1]][1], day2[[1]][1], day3[[1]][1], day4[[1]][1], day5[[1]][1],
               day6[[1]][1], day7[[1]][1])
plot(incr.conv, type="h", xlab="Days after Email Sent", ylab="Increase in Conversion Rate")
abline(h=0)
```

Ad response over time for 2017-01-03 email test



In the test on 2017-01-03, the lift in conversion rate falls to zero in about three days.

Holdout testing

Planning a holdout test

Designing a holdout test

How do you measure response? For how long?

- Conversion (binary), Sales (number), or Site Visits (count)
- During or after campaign

What advertising do the treatment and control groups receive?

- How many ads does the treatment group see? Usual? Heavy up?
- What is used for the control? “Go dark”? alternative ad? Usual?

Which customers are included in the test?

- Current customers, good customers or prospects?

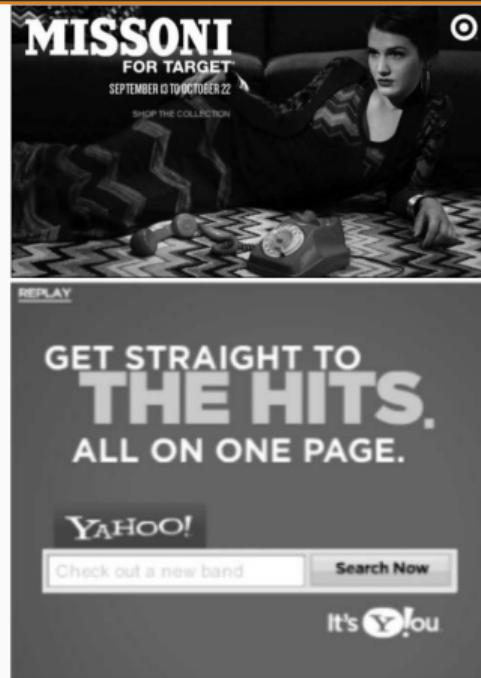
How should customers be assigned to treatment and control groups?

- Always random, but doesn't have to be an equal split.

Holdout testing

**Example: Johnson, Lewis and Reiley
(2017)**

Display advertising holdout test at Yahoo!



During two weeks in Spring 2010, a national apparel retailer ran a test on Yahoo! to measure response to **display ads** on Yahoo!.

The ads displayed the retailer brand and brands of apparel firm, with slideshow transitions between photographs and text.

This test is described in detail in detail in [Johnson, Lewis and Reiley \(2017\)](#) "When less is more: Data and power in advertising experiments", *Marketing Science* 36(1), 43-53.

Source: [Johnson, Lewis and Reiley \(2017\)](#)

Display advertising holdout test: design

How do you measure response? For how long?

- Sales (\$) at the target retailer during the two week campaign

Display advertising holdout test: design

How do you measure response? For how long?

- Sales (\$) at the target retailer during the two week campaign

What advertising do the treatment and control groups receive?

- **Full**: all retailer ads at all opportunities
- **Control**: house Yahoo! ads at all opportunities
- **Half**: half retailer and half house ads

Display advertising holdout test: design

How do you measure response? For how long?

- Sales (\$) at the target retailer during the two week campaign

What advertising do the treatment and control groups receive?

- **Full**: all retailer ads at all opportunities
- **Control**: house Yahoo! ads at all opportunities
- **Half**: half retailer and half house ads

Which customers are included in the test?

- Registered Yahoo! users who were also customers of the retailer were assigned to groups at the start of the two week period. Many of these users never visit Yahoo! during the campaign and don't have the opportunity to see the ads.
- Database match allows us to track users between Yahoo! and the retailer.

Display advertising holdout test: design

How do you measure response? For how long?

- Sales (\$) at the target retailer during the two week campaign

What advertising do the treatment and control groups receive?

- **Full**: all retailer ads at all opportunities
- **Control**: house Yahoo! ads at all opportunities
- **Half**: half retailer and half house ads

Which customers are included in the test?

- Registered Yahoo! users who were also customers of the retailer were assigned to groups at the start of the two week period. Many of these users never visit Yahoo! during the campaign and don't have the opportunity to see the ads.
- Database match allows us to track users between Yahoo! and the retailer.

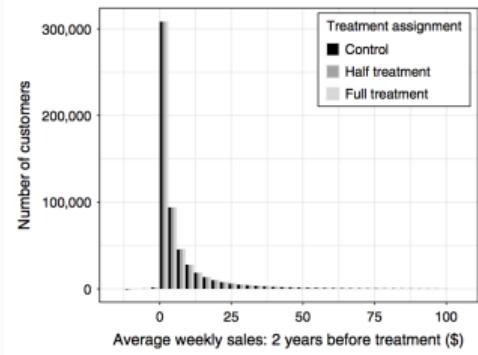
How should customers be assigned to treatment and control groups?

- Randomly, with one third in each group

Display advertising holdout test: randomization checks

It is a good idea to check that the treatment and control groups looked similar prior to the test.

	Treatment group			p-value
	Full	Half	Control	
Sample size	1,032,204	1,032,074	1,032,299	0.988
Female (mean)	68.5%	68.5%	68.5%	0.794
Age (mean)	43.6	43.6	43.6	0.607
Yahoo! page views ^a (mean)	245.8	244.4	243.5	0.132 ^d
Pretreatment sales (two years, mean)	\$857.74	\$859.30	\$855.54	0.475
Pretreatment sales (two weeks, mean)	\$19.34	\$19.24	\$19.10	0.517
Treated subsample				
Exposed sample	572,574	571,222	570,908	0.254
Yahoo! page views (mean)	412.2	411.5	410.1	0.108
Ad views (mean)	33.42	33.41	33.66	0.164
Ad views (median)	15	15	15	
Retailer ad views (mean)	33.42	16.69	—	0.801
Control ad views (mean)	—	16.72	33.66	0.165
Retailer ad click-through rate ^b (%)	0.19	0.24	—	
Retailer ad clicker rate ^c (%)	4.91	3.39	—	



Source: [Johnson, Lewis and Reiley \(2017\)](#)

The three groups appear to be similar. So, the randomization worked and we are comparing “apples to apples”.

Display advertising holdout test: result

Sales Increase for Full and Half Treatments (\$)

	(1)	(2)
Subset of users ^a	Everyone	Treated
Sales after first ad exposure ^b		
Full treatment (\$)	0.673** (0.317)	0.525** (0.237)
Half treatment (\$)	0.0248 (0.311)	0.189 (0.235)
Constant (\$)	15.52*** (0.122)	15.53*** (0.166)

Looking at all the customers, the treatment increased sales by \$0.67 (95% CI = [0.05, 1.29]). Among those who actually were exposed to ads, the effect was \$0.52 (95% CI = [0.06, 0.97]).

Holdout testing

**Example: Gordon, Zettelmeyer,
Bahargava and Chapsky (2016WP)**

Social advertising holdout test

During two weeks in the first half of 2015, an omni-channel retailer conducted a test to determine the lift Facebook ads in the newsfeed.

The image displays a Facebook newsfeed interface comparing two groups: 'Test' and 'Control'. On the left, under 'Test', there are four numbered icons corresponding to an 'Ad Auction': 1. A green square with a white stylized 'J' (representing J&J). 2. A blue square with a white building icon (representing Waterford Lux Resorts). 3. A pink square with a white stylized 'L' (representing Luxe Geneva). 4. A green square with a white stylized 'J' (representing J&J). A red arrow points from the first icon to a post by 'J&J' featuring a fig tart recipe. On the right, under 'Control', there are four numbered icons corresponding to an 'Ad Auction': 1. A blue square with a white building icon (representing Waterford Lux Resorts). 2. A pink square with a white stylized 'L' (representing Luxe Geneva). 3. A green square with a white stylized 'J' (representing J&J). 4. A red square with a white stylized 'C' (representing Céline). A red arrow points from the second icon to a post by 'Waterford Lux Resorts' advertising their travel app. Both posts show engagement metrics like likes, comments, and shares.

Source: Gordon et al. (2016 WP)

This test is reported in Gordon, Zettelmeyer, Bhargava and Chapsky (2016WP) "A Comparison of Approaches to Advertising Measurement: Evidence from Big Field Experiments at Facebook"

Search advertising holdout test design

Which customers are included in the test?

- Facebook users that met the targeting criteria for the campaign

Search advertising holdout test design

Which customers are included in the test?

- Facebook users that met the targeting criteria for the campaign

How do you measure response? For how long?

- Purchases at the retailers website measured via Facebook conversion pixel during the campaign period and up to several weeks after the study ended.

Search advertising holdout test design

Which customers are included in the test?

- Facebook users that met the targeting criteria for the campaign

How do you measure response? For how long?

- Purchases at the retailers website measured via Facebook conversion pixel during the campaign period and up to several weeks after the study ended.

How should customers be assigned to treatment and control groups?

- Randomly, 70% test and 30% control

Search advertising holdout test design

Which customers are included in the test?

- Facebook users that met the targeting criteria for the campaign

How do you measure response? For how long?

- Purchases at the retailers website measured via Facebook conversion pixel during the campaign period and up to several weeks after the study ended.

How should customers be assigned to treatment and control groups?

- Randomly, 70% test and 30% control

What advertising do the treatment and control groups receive?

- Test: user is eligible to see ads and will if the retailer wins the Facebook auction
- Control: competitor ads

Social advertising holdout test: Control condition

Since purchasing house ads or public service announcements (PSAs) is expensive, the Facebook experiment ran a competitor ad for the control.

The image displays a Facebook Ad Auction interface. On the left, under 'Test', there is a list of four ads with their corresponding icons:

- 1.
- 2.
- 3.
- 4.

An arrow points from the icon of the green fig to the ad for 'Japon's Market' which features a fig tart with almonds. On the right, under 'Control', there is a list of four ads with their corresponding icons:

- 1.
- 2.
- 3.
- 4.

An arrow points from the icon of the blue building to the ad for 'Waterford Lux Resorts' which features a woman holding a smartphone.

Source: Gordon et al. (2016 WP)

By using the “next advertiser” as the control condition, the sales lift we estimate is relative to “what would have happened had we not done the campaign.”

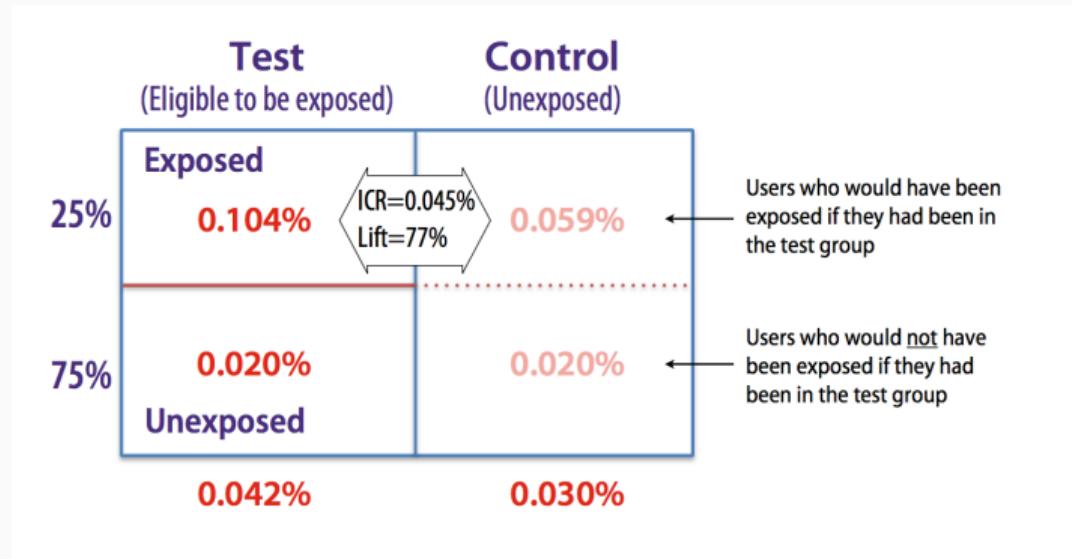
Social holdout test: Randomization check

Variable	Control group	Test group	p-value
Average user age	31.7	31.7	0.33
% of users who are male	17.2%	17.2%	0.705
Length of time using FB (days)	2,288	2,287	0.24
% of users with status “married”	19.6	19.6	0.508
% of users status “engaged”	13.8	13.8	0.0892
% of users status “single”	14.0	14.0	0.888
# of FB friends	485.7	485.7	0.985
# of FB uses in last 7 days	6.377	6.376	0.14
# of FB uses in last 28 days	25.5	25.5	0.172

Source: [Gordon et al. \(2016 WP\)](#)

The test and control groups appear to be similar. Randomization worked.

Social holdout test: Results



Source: Gordon et al. (2016 WP)

There is a 0.045% increase in conversions when comparing the users exposed to ads to the users in the control group that would have been exposed.

Facebook holdout testing tool: Conversion Lift

What is conversion lift?

Conversion lift accurately captures the impact that Facebook ads have in driving business for marketers. Here's how it works:

1. When creating a Facebook campaign, a randomized test group (people that see ads) and control group (people that don't) are established.
2. The advertiser securely shares conversion data from the campaign with Facebook. Typically, this data comes from sources like the Facebook Custom Audiences pixel, conversion pixel or secure point-of-sale (POS) data.
3. Facebook determines additional lift generated from the campaign by comparing conversions in the test and control groups.
4. The results of the study are made available in Ads Manager.

Source: [Facebook Conversion Lift](#)

There were **5,000** additional buyers as a result of your ads, an increase of **3.00%**.



+ 5,000
Incremental Buyers



+ 3.00%
Conversion Lift



GROUP SIZE

CONVERSION RATE

BUYERS

TEST

800K

X

25.0%



200K

CONTROL

200K

4.0x

SCALED
CONTROL

800K

X

24.4%



195K

INCREMENTAL
BUYERS

5,000 ± 1,102

STEP 1: COMPARE THE GROUPS

Scaling the control group to the same size as the test group provides an equal comparison between groups.

STEP 2: APPLY THE CONVERSION RATES

Applying the conversion rate from the control group to the scaled control shows how many people would have bought if the control group were the same size as the test group.

STEP 3: FIND THE INCREMENTAL BUYERS

Subtracting the buyers in the scaled control group from the buyers in the test group gives you the incremental buyers, the number of buyers directly attributable to your ads.

[View Conversion Lift >](#)

[Close this card](#)

Holdout testing

**Example: Blake, Nosko and Tadellis
(2015)**

Search advertising holdout test

Google search results for "used gibson les paul":

Sponsored Ads:

- New: Used Les Paul Gibson**
www.guitarcenter.com/
★★★★★ 12,669 reviews for guitarcenter.com
Free Shipping on 1000's of Items!
2,700 people +1d or follow Guitar Center
\$10 Off \$49 Or \$20 Off \$99+ - Free Shipping to Stores
Special February Financing Locations
- Used Les Paul at Amazon**
www.amazon.com/instruments/
★★★★★ 1,200 seller reviews
Sound Value on Instruments & Gear Over 10,000 Instruments
- Used Gibson Les Paul**
www.ebay.com/
Deals - Used Gibson Les Paul
See NextTag Sellers' Lowest Price!
- Gibson Les Paul Used Sale**
gibson-les-paul-used.compareit.com/
Up To 70% Off Gibson Les Paul Used
Gibson Les Paul Used Compare
- Used Gibson Guitars**
www.williesguitars.com/
Vintage Les Paul, 335, SG, Guitar
Best Prices Fast Shipping & Service
- Win Gibson Les Paul**
bluesmasters.yourjo
Win Gibson Les Paul Guitar
View or Enter Blue Contest
- Gibson Les Paul Used**
www.webcrawler.com/
Search multiple engines for
gibson les paul used
- See your ad here x**

Organic Results:

- Gibson Les Paul Standard**
www.guitarcenter.com/gibson/used/electric-guitar/
25+ items - Welcome to our Gibson Guitars landing page. Dave's Guitar ...
8.6 pounds! \$2,995.00 Gibson '58 Reissue Les Paul Figured Top 12 Ice Tea ...
9.4 pounds! \$2,250.00 Gibson Les Paul Custom Maduro '12
- Gibson | Dave's Guitar Shop**
davesguitar.com/gibson/used/electric-guitar/
25+ items - Welcome to our Gibson Guitars landing page. Dave's Guitar ...
8.6 pounds! \$2,995.00 Gibson '58 Reissue Les Paul Figured Top 12 Ice Tea ...
9.4 pounds! \$2,250.00 Gibson Les Paul Custom Maduro '12
- Gibson Guitar - Get great deals for Gibson Guitar on eBay!**
popular.ebay.com / Popular Items / Musical Instruments
1968 Vintage Gibson Les Paul Standard Gold Top all original. 1 bid. US \$5,000.00 ...
2008 Gibson Les Paul Studio Faded Mahogany Brown USA Electric Guitar. 7 bids ...
Used to \$ Clear Preferences. Buying formats. Auction. Buy It Now ...
- Gibson Les Paul - eBay - Find Popular Products on eBay!**
popular.ebay.com / Popular Items / Musical Instruments
Manufactured by Gibson, the Gibson Les Paul is one of the most widely known electric guitars. ... USED Gibson Les Paul LP Traditional Plus Top Iced Tea ...

In April-July 2012, eBay ran a holdout experiment to determine the sales lift of their search advertising targeting keywords other than "eBay".

This test is reported in detail in [Blake, Nosko and Tadelis \(2015\) "Consumer Heterogeneity and Paid Search Effectiveness: A Large-Scale Field Experiment", *Econometrica* 83\(1\), 155-174.](#)

Source: [Blake, Nosko and Tadelis \(2015\)](#)

Search advertising holdout test: Design

Which customers are included in the test?

- Google AdWords does not provide a testing tool that will randomize search ad exposure at the user level. Instead, they conducted a **geo-test** with 30% of the DMAs in the US market.

Search advertising holdout test: Design

Which customers are included in the test?

- Google AdWords does not provide a testing tool that will randomize search ad exposure at the user level. Instead, they conducted a **geo-test** with 30% of the DMAs in the US market.

How do you measure response? For how long?

- eBay sales (\$) shipped to the DMA

Search advertising holdout test: Design

Which customers are included in the test?

- Google AdWords does not provide a testing tool that will randomize search ad exposure at the user level. Instead, they conducted a **geo-test** with 30% of the DMAs in the US market.

How do you measure response? For how long?

- eBay sales (\$) shipped to the DMA

What advertising do the treatment and control groups receive?

- Test: keep search advertising at usual level
- Control: “go dark”

Search advertising holdout test: Design

Which customers are included in the test?

- Google AdWords does not provide a testing tool that will randomize search ad exposure at the user level. Instead, they conducted a **geo-test** with 30% of the DMAs in the US market.

How do you measure response? For how long?

- eBay sales (\$) shipped to the DMA

What advertising do the treatment and control groups receive?

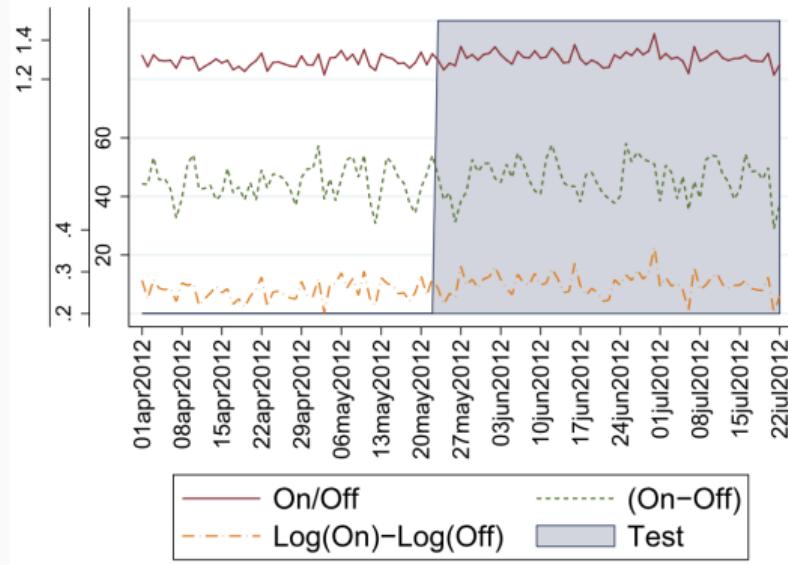
- Test: keep search advertising at usual level
- Control: “go dark”

How should DMAs be assigned to treatment and control groups?

- DMAs were **matched** based on prior eBay sales and then randomly assigned to groups.

Search advertising holdout test: Findings

Ratio of Sales
Before and During Test Period



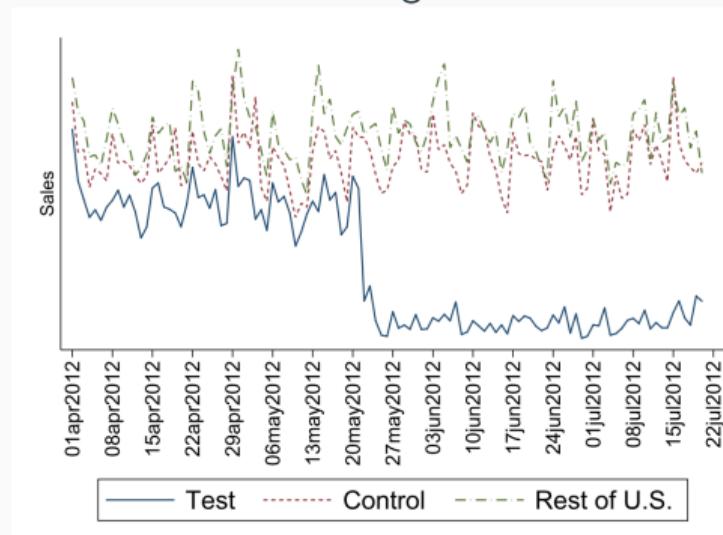
Search advertising does not increase eBay sales.
As a result, eBay has drastically reduced search advertising.

Source: [Blake, Nosko and Tadellis \(2015\)](#)

Search advertising holdout test: “Last Touch”

“Last Touch” Attributed Sales
(sale within 24 hours of search ad exposure)
Before and During Test Period

If you were using last touch attribution, you would think that search advertising made a big difference.



Source: [Blake, Nosko and Tadellis \(2015\)](#)

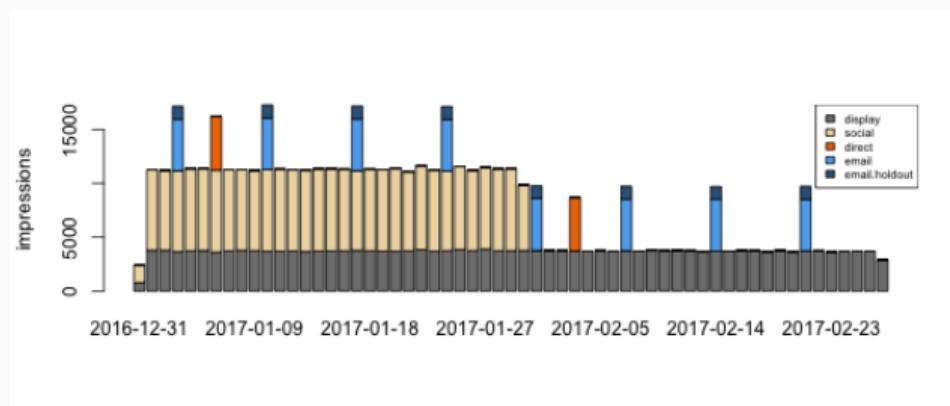
Marketing mix modeling

Marketing mix modeling

What is marketing mix modeling?

Marketing mix modeling

Before advertisers had access to user-level ad exposure data (like our example data set) they only had data on what they had spent on advertising and (sometimes) the number of people who saw the ad. This type of spending data is similar to the aggregate impressions data we looked at for our data set.



To find the correlation between **total sales** in each day/week/month to **advertising spending or impressions** on that same day/week/month, you estimate **regression** model.

A simple marketing mix model

A regression is an equation relating a response ("dependent variable") to one or more other variables ("independent variables").

A simple marketing mix model

A regression is an equation relating a response ("dependent variable") to one or more other variables ("independent variables").

A simple marketing mix model is:

$$\text{sales}_t = \beta_0 + \beta_1 \text{display}_t + \beta_2 \text{social}_t + \beta_3 \text{email}_t + \beta_4 \text{direct}_t + \epsilon_t$$

A simple marketing mix model

A regression is an equation relating a response ("dependent variable") to one or more other variables ("independent variables").

A simple marketing mix model is:

$$\text{sales}_t = \beta_0 + \beta_1 \text{display}_t + \beta_2 \text{social}_t + \beta_3 \text{email}_t + \beta_4 \text{direct}_t + \epsilon_t$$

The words represent data we have for each time period and the β 's represent the unknown relationship between ad impressions and sales. For example β_1 is the increase in sales we get for each additional display impression.

A simple marketing mix model

A regression is an equation relating a response ("dependent variable") to one or more other variables ("independent variables").

A simple marketing mix model is:

$$\text{sales}_t = \beta_0 + \beta_1 \text{display}_t + \beta_2 \text{social}_t + \beta_3 \text{email}_t + \beta_4 \text{direct}_t + \epsilon_t$$

The words represent data we have for each time period and the β 's represent the unknown relationship between ad impressions and sales. For example β_1 is the increase in sales we get for each additional display impression.

Using observations of the sales and advertising data, when can estimate these unknown parameters of the model.

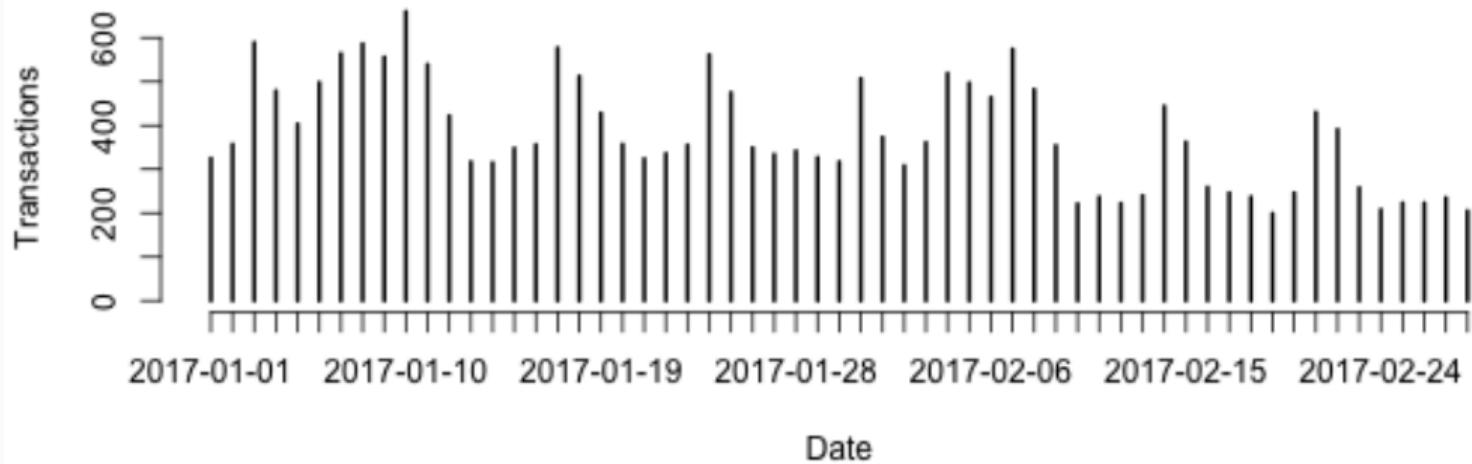
Marketing mix modeling

Aggregating the user-level sales data

As a reminder, we already computed the aggregate transactions for each day from our raw user-level data using a crosstab.

```
> (transbyday <- xtabs(~date, data=trans))  
date  
2017-01-01 2017-01-02 2017-01-03 2017-01-04 2017-01-05 2017-01-06 2017-01-07  
      325        357        589        479        403        498        564  
2017-01-08 2017-01-09 2017-01-10 2017-01-11 2017-01-12 2017-01-13 2017-01-14  
      586        556        660        539        422        317        315  
2017-01-15 2017-01-16 2017-01-17 2017-01-18 2017-01-19 2017-01-20 2017-01-21  
      348        356        577        512        428        357        324
```

Sales by day



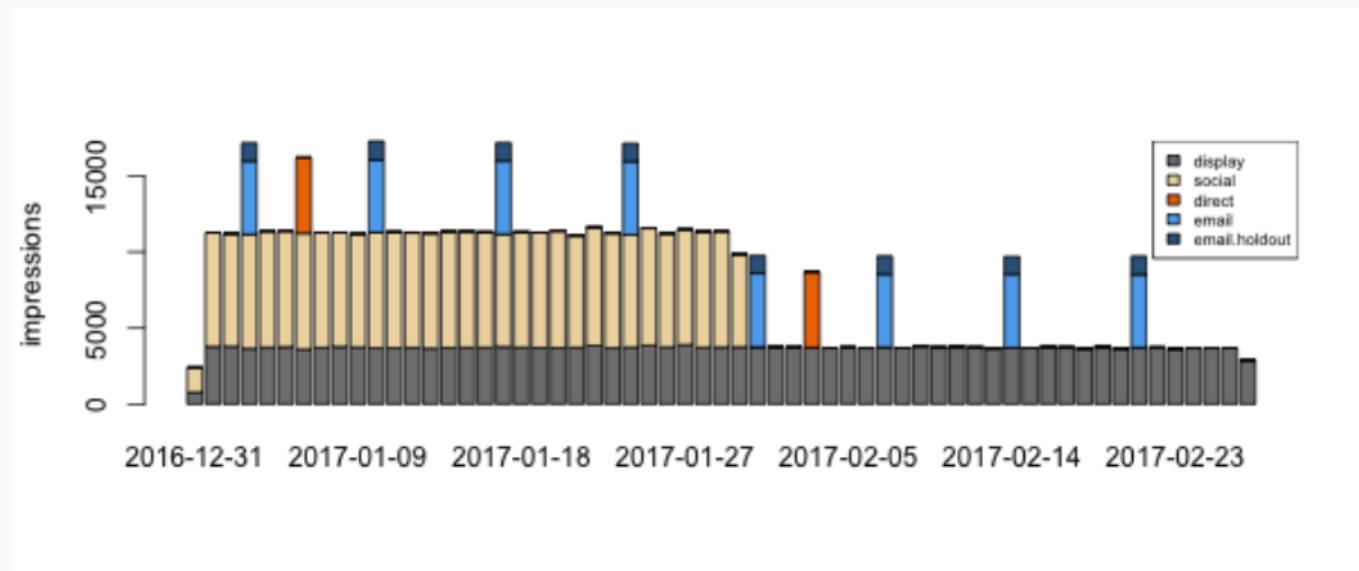
Aggregating the user-level advertising data

We also computed the aggregate number of impressions on each day.

```
> (cadence <- xtabs(~date+channel, data=impress))
```

	channel	direct	display	email	email.holdout	social
date						
2016-12-31		0	788	0	0	1610
2017-01-01		0	3786	0	0	7481
2017-01-02		0	3792	0	0	7416
2017-01-03		0	3656	4798	1203	7505
2017-01-04		0	3731	0	0	7648
2017-01-05		0	3770	0	0	7620
2017-01-06	4974	3611	0	0	0	7614
2017-01-07		0	3719	0	0	7552
2017-01-08		0	3780	0	0	7504
2017-01-09		0	3744	0	0	7446

Advertising by day



Combining the sales and impressions data

When we fit models, it will be convenient to put the sales and impressions data together in the same data frame.

```
> mdata <- as.data.frame(cbind(trans=transbyday[1:57], cadence[2:58,])) # aligning  
> head(mdata)
```

	trans	direct	display	email	email.holdout	social
2017-01-01	325	0	3786	0	0	7481
2017-01-02	357	0	3792	0	0	7416
2017-01-03	589	0	3656	4798	1203	7505
2017-01-04	479	0	3731	0	0	7648
2017-01-05	403	0	3770	0	0	7620
2017-01-06	498	4974	3611	0	0	7614

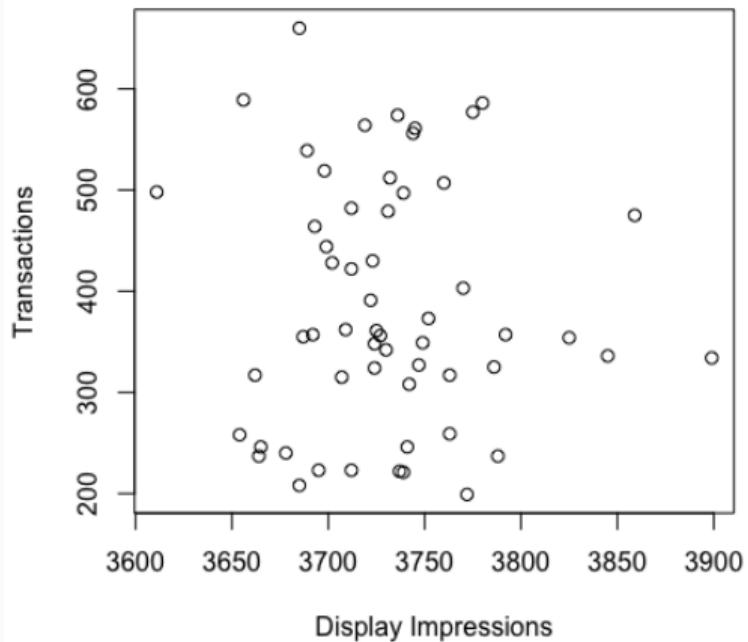
Data exploration

Now that we have the data aggregated, we should do some exploration before modeling.

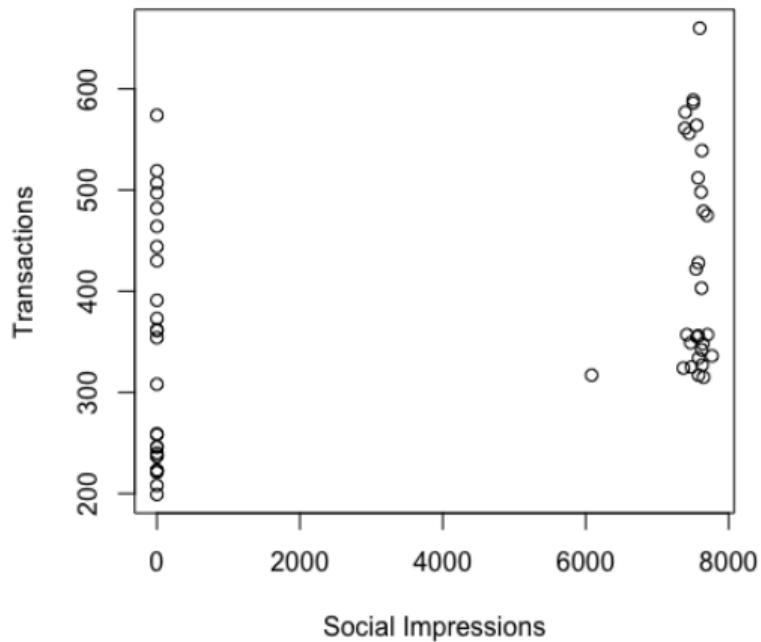
```
> plot(x=mdata$direct, y=mdata$trans, xlab="Direct Impressions",
+       ylab="Transactions", main="Direct Impressions v. Transactions")
> plot(x=mdata$email, y=mdata$trans, xlab="Email Impressions",
+       ylab="Transactions", main="Email Impressions v. Transactions")
> plot(x=mdata$display, y=mdata$trans, xlab="Display Impressions",
+       ylab="Transactions", main="Display Impressions v. Transactions")
> plot(x=mdata$social, y=mdata$trans, xlab="Social Impressions",
+       ylab="Transactions", main="Social Impressions v. Transactions")
```

Daily display and social impressions versus transactions

Display Impressions v. Transactions

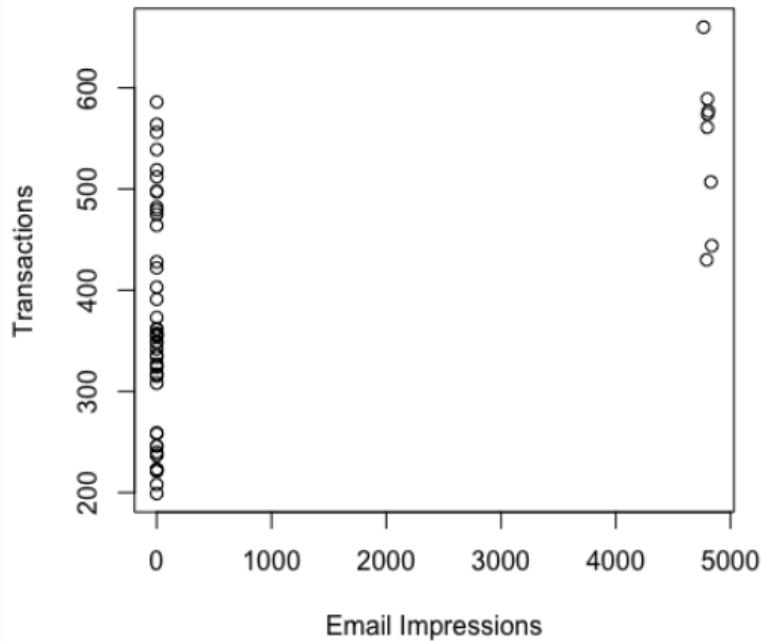


Social Impressions v. Transactions

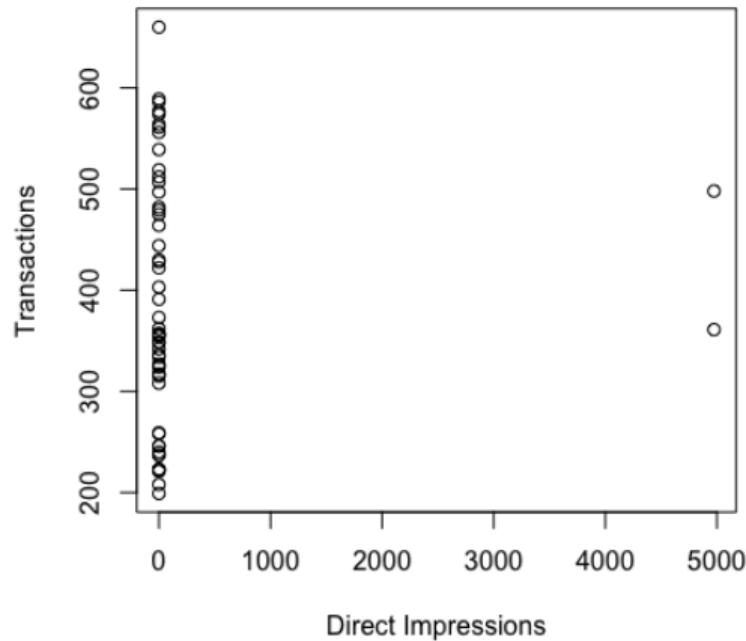


Daily email and direct mpressions versus transactions

Email Impressions v. Transactions



Direct Impressions v. Transactions



Running a simple media mix model

```
> m1 <- lm(trans~direct+social+email+display+social+email.holdout, data=mdata)
> summary(m1)
...
Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) 438.501429  940.622551   0.466 0.643070
direct       0.014529   0.013853   1.049 0.299206
social       0.012663   0.003379   3.747 0.000457 ***
email        -0.161643   0.321735  -0.502 0.617541
display      -0.035299   0.252371  -0.140 0.889313
email.holdout  0.805475   1.290198   0.624 0.535211
---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 92.53 on 51 degrees of freedom
Multiple R-squared:  0.4644,    Adjusted R-squared:  0.4119
F-statistic: 8.845 on 5 and 51 DF,  p-value: 0.00000427
```

Things to notice in the model output

```
> m1 <- lm(trans~direct+social+email+display+social+email.holdout,
   data=mdata)
> summary(m1)

Call:
lm(formula = trans ~ direct + social + email + display + social +
   email.holdout, data = mdata)

Residuals:
    Min      1Q  Median      3Q     Max 
-106.35 -67.79 -41.56  54.42 211.04 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) 438.501429  940.622551   0.466  0.643070    
direct       0.014529   0.013853    1.049  0.299206    
social       0.012663   0.003379    3.747  0.000457 ***  
email        -0.161643   0.321735   -0.502  0.617541    
display      -0.035299   0.252371   -0.140  0.889313    
email.holdout 0.805475   1.290198    0.624  0.535211    
---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 92.53 on 51 degrees of freedom
Multiple R-squared:  0.4644,    Adjusted R-squared:  0.4119 
F-statistic: 8.845 on 5 and 51 DF,  p-value: 0.00000427
```

The only statistically significant effect is for social impressions, and we get 0.012663 additional transactions for each social impression.

Things to notice in the model output

```
> m1 <- lm(trans~direct+social+email+display+social+email.holdout,
   data=mdata)
> summary(m1)

Call:
lm(formula = trans ~ direct + social + email + display + social +
    email.holdout, data = mdata)

Residuals:
    Min      1Q  Median      3Q     Max 
-106.35 -67.79 -41.56  54.42 211.04 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) 438.501429  940.622551   0.466  0.643070    
direct       0.014529   0.013853   1.049  0.299206    
social       0.012663   0.003379   3.747  0.000457 ***  
email        -0.161643   0.321735  -0.502  0.617541    
display      -0.035299   0.252371  -0.140  0.889313    
email.holdout 0.805475   1.290198   0.624  0.535211    
---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 92.53 on 51 degrees of freedom
Multiple R-squared:  0.4644,    Adjusted R-squared:  0.4119 
F-statistic: 8.845 on 5 and 51 DF,  p-value: 0.00000427
```

The only statistically significant effect is for social impressions, and we get 0.012663 additional transactions for each social impression.

The standard error for display is very large, which means we don't have a precise estimate of the effect of display. This happened because daily display impressions are pretty much the same every day. The data is not informative!

Things to notice in the model output

```
> m1 <- lm(trans~direct+social+email+display+social+email.holdout,
   data=mdata)
> summary(m1)

Call:
lm(formula = trans ~ direct + social + email + display + social +
    email.holdout, data = mdata)

Residuals:
    Min      1Q  Median      3Q     Max 
-106.35 -67.79 -41.56  54.42 211.04 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) 438.501429  940.622551   0.466  0.643070    
direct       0.014529   0.013853   1.049  0.299206    
social       0.012663   0.003379   3.747  0.000457 ***  
email        -0.161643   0.321735  -0.502  0.617541    
display      -0.035299   0.252371  -0.140  0.889313    
email.holdout 0.805475   1.290198   0.624  0.535211    
---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 92.53 on 51 degrees of freedom
Multiple R-squared:  0.4644,    Adjusted R-squared:  0.4119 
F-statistic: 8.845 on 5 and 51 DF,  p-value: 0.00000427
```

The only statistically significant effect is for social impressions, and we get 0.012663 additional transactions for each social impression.

The standard error for display is very large, which means we don't have a precise estimate of the effect of display. This happened because daily display impressions are pretty much the same every day. The data is not informative!

The estimated effect of email and display is negative (but not significant).

Adding other variables

In addition to advertising impressions, we might want to include other predictors of transactions in our model, such as the day of the week.

```
> mdata$dayofweek <- weekdays(as.Date(rownames(mdata)))
> m2 <- lm(trans ~ email + direct + display + social + dayofweek, data = mdata)
```

Our second model

```
> summary(m2)

Call:
lm(formula = trans ~ email + direct + display + social + dayofweek,
    data = mdata)

Residuals:
    Min      1Q  Median      3Q     Max 
-87.09 -51.78 -21.70  25.46 221.46 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) 348.192012  891.949059   0.390 0.698063  
email        -0.960694   1.463482  -0.656 0.514809  
direct        0.029537   0.014165   2.085 0.042632 *  
display       -0.031124   0.239069  -0.130 0.896986  
social         0.012754   0.003077   4.144 0.000145 *** 
dayofweekMonday 71.405198  45.616341   1.565 0.124357  
dayofweekSaturday 64.445464  45.601326   1.413 0.164318  
dayofweekSunday  55.293919  44.598357   1.240 0.221330  
dayofweekThursday 67.663736  45.341037   1.492 0.142440  
dayofweekTuesday 4877.229334 7029.997193  0.694 0.491313  
dayofweekWednesday 171.086874  45.336050   3.774 0.000459 *** 
...
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 83.94 on 46 degrees of freedom
Multiple R-squared:  0.6024,    Adjusted R-squared:  0.516 
F-statistic:  6.97 on 10 and 46 DF,  p-value: 0.000001582
```

When we add other variables, it can change the coefficients of the model.

Our second model

```
> summary(m2)

Call:
lm(formula = trans ~ email + direct + display + social + dayofweek,
    data = mdata)

Residuals:
    Min      1Q  Median      3Q     Max 
-87.09 -51.78 -21.70  25.46 221.46 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) 348.192012  891.949059   0.390 0.698063    
email        -0.960694   1.463482  -0.656 0.514809    
direct        0.029537   0.014165   2.085 0.042632 *  
display       -0.031124   0.239069  -0.130 0.896986    
social         0.012754   0.003077   4.144 0.000145 *** 
dayofweekMonday 71.405198  45.616341   1.565 0.124357    
dayofweekSaturday 64.445464  45.601326   1.413 0.164318    
dayofweekSunday  55.293919  44.598357   1.240 0.221330    
dayofweekThursday 67.663736  45.341037   1.492 0.142440    
dayofweekTuesday 4877.229334 7029.997193   0.694 0.491313    
dayofweekWednesday 171.086874  45.336050   3.774 0.000459 *** 
...
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 83.94 on 46 degrees of freedom
Multiple R-squared:  0.6024,    Adjusted R-squared:  0.516 
F-statistic:  6.97 on 10 and 46 DF,  p-value: 0.000001582
```

When we add other variables, it can change the coefficients of the model.

With our new model we find a significant association between direct impressions and transactions (29.5 additional transactions / 1000 direct impressions) in addition to social (12.8 additional transactions / 1000 impressions).

Summary: marketing mix models

A **marketing mix model** is a **regression** relating advertising spending or total impressions to some response like sales.

To build a marketing mix model, you assemble data on advertising spending or impressions by day/week/month and conversions or sales by day/week/month and then fit a model using any regression tool.

The coefficients of this model tell us the correlations between daily/weekly/monthly levels of advertising and sales.

What can go wrong with marketing mix modeling?

There isn't enough variation in the advertising over time or two different types of advertising tend to happen at the same time.

- Model estimates will show large standard errors on the coefficients. You'll be stuck, but at least you'll know it.

What can go wrong with marketing mix modeling?

There isn't enough variation in the advertising over time or two different types of advertising tend to happen at the same time.

- Model estimates will show large standard errors on the coefficients. You'll be stuck, but at least you'll know it.

Advertising happens during periods of peak demand, reversing the causality.

- Model estimates may appear reasonable but will overestimate advertising response. You'll be wrong, but you won't know it.

What can go wrong with marketing mix modeling?

There isn't enough variation in the advertising over time or two different types of advertising tend to happen at the same time.

- Model estimates will show large standard errors on the coefficients. You'll be stuck, but at least you'll know it.

Advertising happens during periods of peak demand, reversing the causality.

- Model estimates may appear reasonable but will overestimate advertising response. You'll be wrong, but you won't know it.

The analyst can twittle with the model until she gets it to say what she likes or what you want to hear.

Model-based attribution

Model-based attribution

What is model-based attribution?

Model-based attribution

Model-based attribution is similar to marketing mix modeling, but we do the analysis at the user-level, relating a user's transactions or conversions to her prior advertising exposures. A simple attribution model is:

$$\text{conversion}_{it} = \beta_0 + \beta_1 \text{display}_{it} + \beta_2 \text{social}_{it} + \beta_3 \text{email}_{it} + \beta_4 \text{direct}_{it} + \epsilon_{it}$$

The key difference from a marketing mix model is that now all the variables are indexed by i in addition to t .

Model-based attribution

Model-based attribution

Model-based attribution

To run this regression, we need to create a data frame where each row is a user-day and summarizes the user's impressions and transactions on that day. We have to transform the raw data to this:

```
> head(adata)
  id      date direct display email email.holdout social trans past.purchase has.email has.direct
1 1 2017-01-01     0      0          0      1      0          0      0          0      1
2 1 2017-01-02     0      0          0      1      0          0      0          0      1
3 1 2017-01-03     0      0          0      0      0          0      0          0      1
4 1 2017-01-04     0      0          0      0      0          0      0          0      1
5 1 2017-01-05     0      0          0      1      0          0      0          0      1
6 1 2017-01-06     1      0          0      0      2          0      0          0      1
```

My R code for doing this is a bit messy and takes several steps. I'm sure you could do better with [tidyverse](#).

Data transform step 1: Summary of impressions by user and day

First, we create a data frame that has a daily summary of each user's transactions.

```
> adatal <- as.data.frame(xtabs(~ id + date + channel, data=impress), stringsAsFactors=FALSE)
> adatal$id <- as.integer(adatal$id)
> adatal$date <- as.Date(adatal$date)
> adatal$channel <- as.factor(adatal$channel)
> dimnames(adatal)[[2]][4] <- "impr"
> head(adatal)
   id      date channel impr
1  1 2016-12-31  direct    0
2  2 2016-12-31  direct    0
3  3 2016-12-31  direct    0
4  5 2016-12-31  direct    0
5  6 2016-12-31  direct    0
6  7 2016-12-31  direct    0
```

Data transform step 2: Adding in users with zero impressions

Because some users don't appear in the impressions file, we have to add in rows (with zeros) for those users.

```
> pop <- unique(cust$id)
> no.impress.ids <- pop[!(pop %in% unique(impress$id))]
> dates <- sort(unique(impress$date))
> channels <- unique(impress$channel)
> no.impress.obs <- data.frame(id=rep(no.impress.ids, each=length(dates)*length(channels)),
+                                 date=rep(rep(dates, each=length(channels)), length(no.impress.ids)),
+                                 channel=rep(channels, length(no.impress.ids)*length(dates)),
+                                 impr=rep(0, length(dates)*length(no.impress.ids)*length(channels)),
+                                 stringsAsFactors=FALSE)
> no.impress.obs$channel <- as.factor(no.impress.obs$channel)
> adatal <- rbind(adatal, no.impress.obs)
> summary(adatal)
      id           date          channel        impr
Min. : 1   Min. :2016-12-31   direct     :590000  Min. : 0.0000
1st Qu.: 2501 1st Qu.:2017-01-14   display    :590000  1st Qu.: 0.0000
Median : 5000 Median :2017-01-29   email      :590000  Median : 0.0000
Mean   : 5000 Mean  :2017-01-29   email.holdout:590000  Mean   : 0.1699
3rd Qu.: 7500 3rd Qu.:2017-02-13   social     :590000  3rd Qu.: 0.0000
Max.   :10000  Max.  :2017-02-27                    Max.  :16.0000
```

Data transform step 3: Switch to wide format

Next, we unstack the impressions column by the channel.

```
> adata <- reshape(adata[, direction="wide", v.names="impr", idvar=c("id", "date"),
+                     timevar="channel", new.row.names=NULL)
> sum(adata$impr.direct) == length(impress$channel[impress$channel=="direct"]) #quick check
[1] TRUE
> nrow(adata)
[1] 590000
> head(adata)
   id      date impr.direct impr.display impr.email impr.email.holdout impr.social
1  1 2016-12-31          0          0          0                  0          0
2  2 2016-12-31          0          0          0                  0          0
3  3 2016-12-31          0          0          0                  0          0
4  5 2016-12-31          0          1          0                  0          0
5  6 2016-12-31          0          0          0                  0          0
6  7 2016-12-31          0          0          0                  0          0
```

Data transformation step 4: Add the daily transactions for each user

Finally, we summarize the transactions by customer-day and merge with impressions.

```
> atrans <- as.data.frame(xtabs(~ id + date, data=trans), stringsAsFactors=FALSE)
> atrans$id <- as.integer(atrans$id)
> atrans$date <- as.Date(atrans$date)
> dimnames(atrans)[[2]][3] <- "trans"
> adata <- merge(adata, atrans, by=c("id", "date"), all=TRUE)
> adata$trans[is.na(adata$trans)] <- 0 # fill in zeros for transactions
> head(adata)
```

	id	date	impr.direct	impr.display	impr.email	impr.email.holdout	impr.social	trans
1	1	2016-12-31	0	0	0	0	0	0
2	1	2017-01-01	0	0	0	0	1	0
3	1	2017-01-02	0	0	0	0	1	0
4	1	2017-01-03	0	0	0	0	0	0
5	1	2017-01-04	0	0	0	0	0	0
6	1	2017-01-05	0	0	0	0	1	0

Data transformation step 5: Final tidy up of attribution modeling data

```
> # Remove first and last days (which are incomplete)
> adata <- adata[adata$date!="2016-12-31" & adata$date != "2017-02-28" & adata$date != "2017-02-27",]
# Add customer info from cust table
> adata <- merge(adata, cust, by=c("id"))
# Tidy up column names
> dimnames(adata)[[2]][3:11] <- c("direct", "display", "email", "email.holdout", "social",
   "trans", "past.purchase", "has.email", "has.direct")
> rm(adatal, atrans)
```

Summary of attribution modeling data

```
> summary(adata)
      id          date        direct       display       email
Min.   : 1   Min.   :2017-01-01   Min.   :0.00000   Min.   : 0.0000   Min.   :0.00000
1st Qu.: 2501 1st Qu.:2017-01-15  1st Qu.:0.00000  1st Qu.: 0.0000  1st Qu.:0.00000
Median : 5000 Median :2017-01-29  Median :0.00000  Median : 0.0000  Median :0.00000
Mean   : 5000 Mean   :2017-01-29  Mean   :0.01745  Mean   : 0.3731  Mean   :0.06741
3rd Qu.: 7500 3rd Qu.:2017-02-12  3rd Qu.:0.00000  3rd Qu.: 0.0000  3rd Qu.:0.00000
Max.   :10000 Max.   :2017-02-26  Max.   :1.00000  Max.   :13.0000  Max.   :1.00000
email.holdout    social        trans      past.purchase has.email
Min.   :0.00000   Min.   : 0.0000   Min.   :0.00000   Min.   :0.0000   Min.   :0.0000
1st Qu.:0.00000  1st Qu.: 0.0000  1st Qu.:0.00000  1st Qu.: 0.0000  1st Qu.:0.0000
Median :0.00000  Median : 0.0000  Median :0.00000  Median :1.0000  Median :1.0000
Mean   :0.01681  Mean   : 0.3954  Mean   :0.03858  Mean   : 0.5022  Mean   :0.6001
3rd Qu.:0.00000  3rd Qu.: 0.0000  3rd Qu.:0.00000  3rd Qu.:1.0000  3rd Qu.:1.0000
Max.   :1.00000  Max.   :16.0000  Max.   :1.00000  Max.   :1.0000  Max.   :1.0000
has.direct
Min.   :0.0000
1st Qu.:0.0000
Median :0.0000
Mean   :0.4974
3rd Qu.:1.0000
Max.   :1.0000
```

Summary of attribution modeling data on a specific day

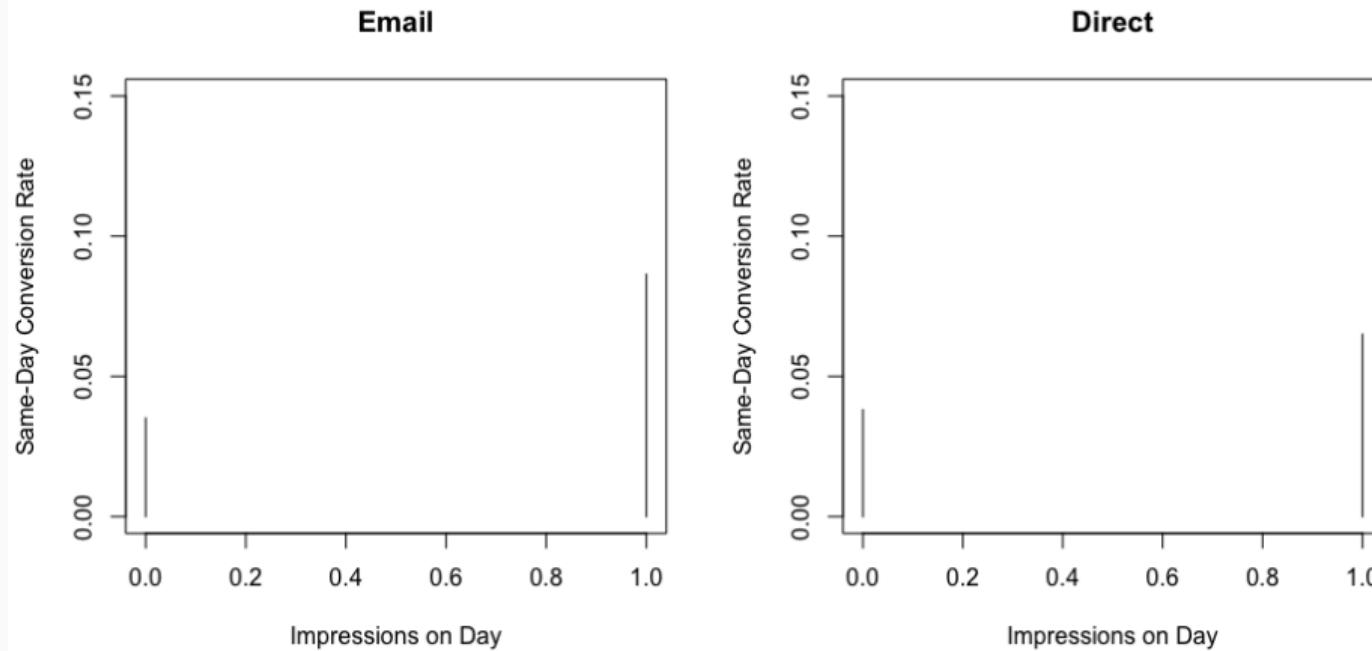
```
> summary(adata[adata$date=="2017-01-03",])
      id          date        direct      display       email   email.holdout
Min. : 1  Min. :2017-01-03  Min. :0  Min. :0.0000  Min. :0.0000  Min. :0.0000
1st Qu.: 2501  1st Qu.:2017-01-03  1st Qu.:0  1st Qu.:0.0000  1st Qu.:0.0000  1st Qu.:0.0000
Median : 5000  Median :2017-01-03  Median :0  Median :0.0000  Median :0.0000  Median :0.0000
Mean   : 5000  Mean   :2017-01-03  Mean   :0  Mean   :0.3656  Mean   :0.4798  Mean   :0.1203
3rd Qu.: 7500  3rd Qu.:2017-01-03  3rd Qu.:0  3rd Qu.:0.0000  3rd Qu.:1.0000  3rd Qu.:0.0000
Max.   :10000  Max.   :2017-01-03  Max.   :0  Max.   :9.0000  Max.   :1.0000  Max.   :1.0000

      social        trans    past.purchase has.email     has.direct
Min. : 0.0000  Min. :0.0000  Min. :0.0000  Min. :0.0000  Min. :0.0000
1st Qu.: 0.0000  1st Qu.:0.0000  1st Qu.:0.0000  1st Qu.:0.0000  1st Qu.:0.0000
Median : 0.0000  Median :0.0000  Median :1.0000  Median :1.0000  Median :0.0000
Mean   : 0.7505  Mean   :0.0589  Mean   :0.5022  Mean   :0.6001  Mean   :0.4974
3rd Qu.: 1.0000  3rd Qu.:0.0000  3rd Qu.:1.0000  3rd Qu.:1.0000  3rd Qu.:1.0000
Max.   :12.0000  Max.   :1.0000  Max.   :1.0000  Max.   :1.0000  Max.   :1.0000
```

Visualizing the relationship between impressions and transactions

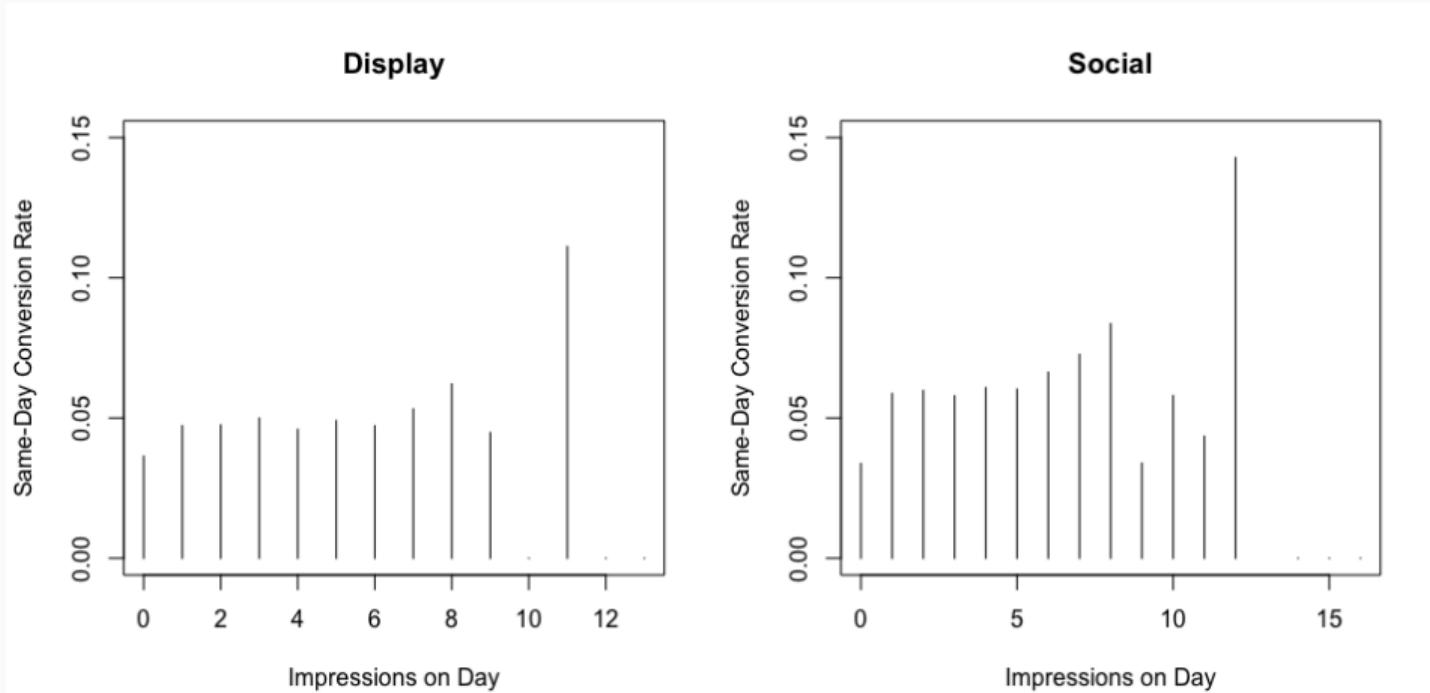
```
> plot(aggregate(trans~direct, data=adata, FUN=mean), type="h", ylim=c(0,0.15),
+       xlab="Impressions on Day", ylab="Same-Day Conversion Rate", main="Direct")
> plot(aggregate(trans~email, data=adata, FUN=mean), type="h", ylim=c(0,0.15),
+       xlab="Impressions on Day", ylab="Same-Day Conversion Rate", main="Email")
> plot(aggregate(trans~display, data=adata, FUN=mean), type="h", ylim=c(0,0.15),
+       xlab="Impressions on Day", ylab="Same-Day Conversion Rate", main="Display")
> plot(aggregate(trans~social, data=adata, FUN=mean), type="h", ylim=c(0,0.15),
+       xlab="Impressions on Day", ylab="Same-Day Conversion Rate", main="Social")
```

Direct and email daily impressions versus transactions



Looking at user-level data, it is easy to see that users convert more on days they get emails or direct mail.

Display and social daily impressions versus transactions



A simple attribution model

```
> m1 <- lm(trans ~ direct + display + email + social, data=adata)
> summary(m1)
...
Coefficients:
            Estimate Std. Error t value      Pr(>|t|)
(Intercept) 0.0296060  0.0003024   97.91 <0.0000000000000002 ***
direct       0.0306548  0.0019423   15.78 <0.0000000000000002 ***
display      0.0041618  0.0002759   15.09 <0.0000000000000002 ***
email        0.0520472  0.0010144   51.31 <0.0000000000000002 ***
social       0.0085516  0.0002542   33.64 <0.0000000000000002 ***
---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 0.1919 on 569995 degrees of freedom
Multiple R-squared:  0.007249,  Adjusted R-squared:  0.007242
F-statistic: 1041 on 4 and 569995 DF,  p-value: < 0.0000000000000022
```

Key things to notice about our attribution model

```
> m1 <- lm(trans ~ direct + display + email + social, data=adata)
> summary(m1)
...
Coefficients:
            Estimate Std. Error t value      Pr(>|t|)
(Intercept) 0.0296060  0.0003024   97.91 <0.0000000000000002 ***
direct       0.0306548  0.0019423   15.78 <0.0000000000000002 ***
display      0.0041618  0.0002759   15.09 <0.0000000000000002 ***
email        0.0520472  0.0010144   51.31 <0.0000000000000002 ***
social       0.0085516  0.0002542   33.64 <0.0000000000000002 ***
---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 0.1919 on 569995 degrees of freedom
Multiple R-squared:  0.007249,  Adjusted R-squared:  0.007242
F-statistic: 1041 on 4 and 569995 DF,
p-value: < 0.000000000000022
```

We have gotten reasonable estimates of advertising response even for this simple model that only looks at same-day response. For example, 1 email impressions is associated with 0.052 additional transactions on the same day.

Key things to notice about our attribution model

```
> m1 <- lm(trans ~ direct + display + email + social, data=adata)
> summary(m1)
...
Coefficients:
            Estimate Std. Error t value      Pr(>|t|)
(Intercept) 0.0296060  0.0003024   97.91 <0.0000000000000002 ***
direct       0.0306548  0.0019423   15.78 <0.0000000000000002 ***
display      0.0041618  0.0002759   15.09 <0.0000000000000002 ***
email        0.0520472  0.0010144   51.31 <0.0000000000000002 ***
social       0.0085516  0.0002542   33.64 <0.0000000000000002 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.1919 on 569995 degrees of freedom
Multiple R-squared:  0.007249, Adjusted R-squared:  0.007242
F-statistic: 1041 on 4 and 569995 DF,
p-value: < 0.000000000000022
```

We have gotten reasonable estimates of advertising response even for this simple model that only looks at same-day response. For example, 1 email impressions is associated with 0.052 additional transactions on the same day.

The probability of transacting on a day where the user is exposed to no advertising is 3%. It would be bad to assume that all sales are due to advertising.

A model that includes user-characteristics

Because we have a user-level model, we can bring in user characteristics.

```
> m2 <- lm(trans ~ direct + display + email + social + past.purchase, data=adata)
> summary(m2)

Coefficients:
              Estimate Std. Error t value     Pr(>|t|)    
(Intercept) 0.0142352  0.0003889   36.60 <0.000000000000002 *** 
direct       0.0203596  0.0019427   10.48 <0.000000000000002 *** 
display      0.0041846  0.0002749   15.22 <0.000000000000002 *** 
email        0.0433566  0.0010205   42.49 <0.000000000000002 *** 
social        0.0086049  0.0002533   33.97 <0.000000000000002 *** 
past.purchase 0.0320725  0.0005130   62.52 <0.000000000000002 *** 
---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 0.1912 on 569994 degrees of freedom
Multiple R-squared:  0.01401,    Adjusted R-squared:  0.014 
F-statistic:  1620 on 5 and 569994 DF,  p-value: < 0.0000000000000022
```

A model that includes user-characteristics

```
> m2 <- lm(trans ~ direct + display + email + social +
>           past.purchase, data=adata)
> summary(m2)

Coefficients:
            Estimate Std. Error t value     Pr(>|t|)
(Intercept) 0.0142352  0.0003889   36.60 <0.0000000000000002 ***
direct       0.0203596  0.0019427   10.48 <0.0000000000000002 ***
display      0.0041846  0.0002749   15.22 <0.0000000000000002 ***
email        0.0433566  0.0010205   42.49 <0.0000000000000002 ***
social       0.0086049  0.0002533   33.97 <0.0000000000000002 ***
past.purchase 0.0320725  0.0005130   62.52 <0.0000000000000002 ***

---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.1912 on 569994 degrees of freedom
Multiple R-squared:  0.01401,    Adjusted R-squared:  0.014
F-statistic: 1620 on 5 and 569994 DF,
p-value: < 0.000000000000022
```

The model suggests that customers who have not made a purchase before are much less likely to make a purchase (1.4% versus 4.6% baseline rate without advertising).

A model that includes user-characteristics

```
> m2 <- lm(trans ~ direct + display + email + social +
>           past.purchase, data=adata)
> summary(m2)

Coefficients:
            Estimate Std. Error t value     Pr(>|t|)
(Intercept) 0.0142352  0.0003889   36.60 <0.0000000000000002 ***
direct       0.0203596  0.0019427   10.48 <0.0000000000000002 ***
display      0.0041846  0.0002749   15.22 <0.0000000000000002 ***
email        0.0433566  0.0010205   42.49 <0.0000000000000002 ***
social       0.0086049  0.0002533   33.97 <0.0000000000000002 ***
past.purchase 0.0320725  0.0005130   62.52 <0.0000000000000002 ***
---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 0.1912 on 569994 degrees of freedom
Multiple R-squared:  0.01401,    Adjusted R-squared:  0.014
F-statistic: 1620 on 5 and 569994 DF,
p-value: < 0.000000000000022
```

The model suggests that customers who have not made a purchase before are much less likely to make a purchase (1.4% versus 4.6% baseline rate without advertising).

The estimates of advertising effects are smaller than in our first model. This is because people who have made a purchase are more likely to be exposed to advertising. When we control for past purchase, we can see that the effects of advertising are smaller than we might have thought based on the first model.

A model that includes user-characteristics

```
> m2 <- lm(trans ~ direct + display + email + social +
>           past.purchase, data=adata)
> summary(m2)

Coefficients:
            Estimate Std. Error t value     Pr(>|t|)
(Intercept) 0.0142352  0.0003889   36.60 <0.0000000000000002 ***
direct       0.0203596  0.0019427   10.48 <0.0000000000000002 ***
display      0.0041846  0.0002749   15.22 <0.0000000000000002 ***
email        0.0433566  0.0010205   42.49 <0.0000000000000002 ***
social       0.0086049  0.0002533   33.97 <0.0000000000000002 ***
past.purchase 0.0320725  0.0005130   62.52 <0.0000000000000002 ***
---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 0.1912 on 569994 degrees of freedom
Multiple R-squared:  0.01401,    Adjusted R-squared:  0.014
F-statistic: 1620 on 5 and 569994 DF,
p-value: < 0.000000000000022
```

The model suggests that customers who have not made a purchase before are much less likely to make a purchase (1.4% versus 4.6% baseline rate without advertising).

The estimates of advertising effects are smaller than in our first model. This is because people who have made a purchase are more likely to be exposed to advertising. When we control for past purchase, we can see that the effects of advertising are smaller than we might have thought based on the first model.

The controls you choose to add to the model can be critical.

Logistic regression

Those of you who learned regression from a careful teacher may have noticed that I was violating a key assumption of the linear regression model.

Logistic regression

Those of you who learned regression from a careful teacher may have noticed that I was violating a key assumption of the linear regression model.

Our outcome variable is binary and so it would be better to use a logistic regression for binary outcomes.

Logistic regression for attribution

```
> m3 <- glm(trans ~ direct + display + email + social + past.purchase,
+             family=binomial(), data=adata)
> summary(m3)

...
Coefficients:
              Estimate Std. Error z value      Pr(>|z|)
(Intercept) -4.022893   0.014013 -287.078 <0.0000000000000002 *** 
direct        0.412554   0.041688    9.896 <0.0000000000000002 *** 
display       0.100630   0.006581   15.291 <0.0000000000000002 *** 
email         0.762355   0.020010   38.099 <0.0000000000000002 *** 
social        0.180079   0.005356   33.621 <0.0000000000000002 *** 
past.purchase 0.957207   0.015672   61.079 <0.0000000000000002 *** 

---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 186297  on 569999  degrees of freedom
Residual deviance: 178869  on 569994  degrees of freedom
AIC: 178881
```

Logistic regression for attribution

```
> m3 <- glm(trans ~ direct + display + email + social + past.purchase,
+             family=binomial(), data=adata)
> summary(m3)

Coefficients:
            Estimate Std. Error z value      Pr(>|z|)
(Intercept) -4.022893   0.014013 -287.078 <0.0000000000000002 ***
direct       0.412554   0.041688    9.896 <0.0000000000000002 ***
display      0.100630   0.006581   15.291 <0.0000000000000002 ***
email        0.762355   0.020010   38.099 <0.0000000000000002 ***
social       0.180079   0.005356   33.621 <0.0000000000000002 ***
past.purchase 0.957207   0.015672   61.079 <0.0000000000000002 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 186297  on 569999  degrees of freedom
Residual deviance: 178869  on 569994  degrees of freedom
AIC: 178881

Number of Fisher Scoring iterations: 6
```

While this model is better, the logistic regression coefficients can be difficult to interpret.

Logistic regression for attribution

```
> m3 <- glm(trans ~ direct + display + email + social + past.purchase,
+             family=binomial(), data=adata)
> summary(m3)

Coefficients:
            Estimate Std. Error z value      Pr(>|z|)
(Intercept) -4.022893   0.014013 -287.078 <0.0000000000000002 ***
direct       0.412554   0.041688    9.896 <0.0000000000000002 ***
display      0.100630   0.006581   15.291 <0.0000000000000002 ***
email        0.762355   0.020010   38.099 <0.0000000000000002 ***
social       0.180079   0.005356   33.621 <0.0000000000000002 ***
past.purchase 0.957207   0.015672   61.079 <0.0000000000000002 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 186297  on 569999  degrees of freedom
Residual deviance: 178869  on 569994  degrees of freedom
AIC: 178881

Number of Fisher Scoring iterations: 6
```

While this model is better, the logistic regression coefficients can be difficult to interpret.

We are still seeing that direct and email have a bigger effect than a single display or social impression.